

Theoretical and empirical runtime analysis of evolutionary algorithms for the PARTITION problem

Bachelor Thesis of

Daniel Lipp

At the Department of Informatics and Mathematics
Chair of algorithms for intelligent systems



Advisor: Prof. Dr. Dirk Sudholt

Time Period: 14th May 2023 – 14th August 2023

Statement of Authorship

I hereby declare that this document has been composed by myself and describes my own work, unless otherwise acknowledged in the text.

Passau, July 28, 2023

Abstract

A short summary of what is going on here.

Deutsche Zusammenfassung

Kurze Inhaltsangabe auf deutsch.

Contents

| | |
|---|-----------|
| 1. Introduction | 1 |
| 2. Preliminaries | 3 |
| 2.1. Notations | 3 |
| 3. Improving bounds on the RLS and (1+1) EA | 5 |
| 3.1. Improving bounds on the RLS and the (1+1) EA | 5 |
| 3.2. Binomial distributed input | 7 |
| 4. Higher mutation rates and heavy tailed mutations | 9 |
| 4.1. Algorithms | 9 |
| 4.2. Runtime Analysis of higher mutation rates | 10 |
| 5. Experimental Results | 13 |
| 5.1. Code | 13 |
| 5.1.1. The Algorithms | 13 |
| 5.1.2. Random number generation | 14 |
| 5.2. Do inputs have perfect partitions? | 15 |
| 5.2.1. Binomial Inputs | 15 |
| 5.3. Binomial distributed values | 18 |
| 5.3.1. RLS Comparison | 20 |
| 5.3.2. (1+1) EA Comparison | 21 |
| 5.3.3. pmut Comparison | 21 |
| 5.3.4. Comparison of the best variants | 22 |
| 5.4. Geometric distributed values | 23 |
| 5.4.1. RLS Comparison | 24 |
| 5.4.2. (1+1) EA Comparison | 24 |
| 5.4.3. pmut Comparison | 24 |
| 5.4.4. Comparison of the best variants | 25 |
| 5.5. Uniform distributed inputs | 26 |
| 5.5.1. RLS Comparison | 26 |
| 5.5.2. (1+1) EA Comparison | 27 |
| 5.5.3. pmut Comparison | 27 |
| 5.5.4. Comparison of the best variants | 28 |
| 5.6. powerlaw distributed inputs | 29 |
| 5.6.1. RLS Comparison | 29 |
| 5.6.2. (1+1) EA Comparison | 30 |
| 5.6.3. pmut Comparison | 31 |
| 5.6.4. Comparison of the best variants | 31 |
| 5.7. OneMax Equivalent for PARTITION | 33 |
| 5.7.1. RLS Comparison | 33 |
| 5.7.2. (1+1) EA Comparison | 34 |

| | | |
|---------------------|---|-----------|
| 5.7.3. | pmut Comparison | 34 |
| 5.7.4. | Comparison of the best variants | 35 |
| 5.8. | Carsten Witt's worst case input | 36 |
| 5.8.1. | RLS Comparison | 37 |
| 5.8.2. | (1+1) EA Comparison | 37 |
| 5.8.3. | pmut Comparison | 38 |
| 5.8.4. | Comparison of the best variants | 38 |
| 5.9. | Multiple distributions mixed | 39 |
| 5.9.1. | RLS Comparison | 39 |
| 5.9.2. | (1+1) EA Comparison | 40 |
| 5.9.3. | pmut Comparison | 40 |
| 5.9.4. | Comparison of the best variants | 40 |
| 5.10. | Adding multiple distributions | 41 |
| 5.10.1. | RLS Comparison | 42 |
| 5.10.2. | (1+1) EA Comparison | 43 |
| 5.10.3. | pmut Comparison | 43 |
| 5.10.4. | Comparison of the best variants | 43 |
| 5.11. | Multiple distributions mixed & overlapped | 44 |
| 5.11.1. | RLS Comparison | 45 |
| 5.11.2. | (1+1) EA Comparison | 45 |
| 5.11.3. | pmut Comparison | 46 |
| 5.11.4. | Comparison of the best variants | 46 |
| 5.12. | Conclusion of empirical results | 47 |
| Bibliography | | 49 |
| Appendix | | 51 |
| A. | Appendix Section 1 | 51 |

1. Introduction

This chapter should contain

1. A short description of the thesis topic and its background.
2. An overview of related work in this field.
3. Contributions of the thesis.
4. Outline of the thesis.

2. Preliminaries

2.1. Notations

1. **RLS**: Randomised Local Search
2. **RSH**: Randomised Search Heuristic referring to all analysed Evolutionary algorithms
3. n : The input length of the problem
4. w_i : The i -th object of the input. If not mentioned otherwise the weights are sorted in non-increasing order so: $w_1 \geq w_2 \geq \dots \geq w_{n-1} \geq w_n$
5. W : The sum of all objects: $W = \sum_{i=1}^n w_i$
6. **bin**: When solving Partition a set of numbers is divided into two distinct subsets and in this paper both subsets are referred to as bins
7. b_F : The fuller bin (the bin with more total weight)
8. b_E : The emptier bin (the bin with less total weight)
9. b_{w_i} : The bin containing the object w_i
10. opt : The optimal solution for a given partition instance.
11. x : A vector $x \in \{0, 1\}^n$ describing a solution

3. Improving bounds on the RLS and (1+1) EA

3.1. Improving bounds on the RLS and the (1+1) EA

Theorem 3.1. *If $w_1 \geq \frac{W}{2}$ then the RLS and the (1+1) EA reach the optimal value in expected time $\Theta(n \log n)$*

Proof. The optimal solution is putting w_1 in one bin and all other elements in the other bin. So the problem is almost identical to OneMax/ZeroMax. A single bit flip of the first bit can only happen, if the emptier bin has a weight of at most $\frac{W-w_1}{2}$. After this flip the weight of the emptier bin is at least $\frac{W-w_1}{2}$ and therefore another single bit flip of w_1 can only happen before a different bit is flipped. The run of the RSHs can be divided into two phases:

Phase 1: The RSH reaches a search point with $b_E \geq \frac{W-w_1}{2}$.

Phase 2: The RSH reaches an optimal solution $\Rightarrow w_1$ is in one bin and all other elements are in the other bin.

The expected length of the first phase is at most en for the (1+1) EA and n for the RLS because the probability of flipping the first bit is at least $\frac{1}{n} \cdot (1 - \frac{1}{n})^{n-1} \geq \frac{1}{en}$ and therefore the expected time for such a step is at most $\frac{1}{\frac{1}{en}} = en$ for the (1+1) EA and $\frac{1}{n} = n$ for the RLS.

In the second phase the RLS can no longer flip w_1 as it does not result in an improvement ever again. Therefore the RLS behaves exactly as on OneMax/ZeroMax depending on the value of the first bit and reaches an optimal solution in $\Theta(n \log n)$ resulting in a total runtime of $\Theta(n \log n)$ (Theorem 3 in [Wit14]). The (1+1) EA has to minimise a linear function of $n - 1$ bits which also takes $\Theta(n \log n)$ time as well (Corollary 4.2 in [Wit13]). Flipping the first bit results in the optimal solution being inverted which could decrease the progress of minimising the linear function. Due to Lemma 4.2 with $\delta = \frac{1}{n}$ (for $n > 1$) the solution is at most $w_1 + \delta(W - w_1)$ after expected time

$$2\lceil en \ln(2/\delta) \rceil = 2\lceil en \ln(2/\frac{1}{n}) \rceil = 2\lceil en \ln(2n) \rceil = 2\lceil en(\ln(n) + \ln(2)) \rceil \leq 2en \ln(n) + 4$$

The value of b_E is then at least $W - (w_1 + \delta(W - w_1)) = (1 - \delta)(W - w_1) = (1 - \frac{1}{n})(W - w_1)$. For a successful flip of w_1 the algorithm must move $2 \cdot (b_E - \frac{W - w_1}{2})$ from b_E to b_F which evaluates to

$$2 \cdot ((1 - \frac{1}{n})(W - w_1) - \frac{W - w_1}{2}) = (W - w_1)(2 - \frac{2}{n} - 1) = (1 - \frac{2}{n})(W - w_1) = \frac{n - 1}{n - 2} b_E$$

TODO: *why can't the first bit be flipped any more? Or why does it not matter?*

So the total expected time is $\mathcal{O}(n) + \Theta(n \log n) = \Theta(n \log n)$ \square

Lemma 3.2. *If $b_F \leq \frac{2}{3} \cdot W$ the approximation ratio is at most $\frac{4}{3}$*

Proof. $\frac{b_F}{opt} \leq \frac{(2/3) \cdot W}{\frac{(2/3) \cdot W}{(1/2) \cdot W}} = \frac{4}{3}$, since $opt \geq \frac{W}{2}$ \square

Corollary 3.3. *If $w_1 \geq \frac{W}{3}$ and w_1 is in the emptier bin, then the approximation ratio is at most $\frac{4}{3}$*

Proof. w_1 is in the emptier bin, so $b_F \leq W - w_1 \leq W - \frac{W}{3} = \frac{2W}{3}$ and with Lemma 3.2 the assumption follows. \square

Lemma 3.4. *Any object of weight v can be moved from b_F to b_E if and only if $b_F - b_E \geq v$*

Proof. " \Leftarrow " :

$b_F - b_E \geq v \Leftrightarrow b_F \geq b_E + v$, so after moving an object with weight v from b_F to b_E , the new weight of b_E is at most the weight of b_F before moving the object, thus the RSH accepts the step.

" \Rightarrow " :

$b_F - b_E < v \Leftrightarrow b_F < b_E + v$, so moving an object of weight v results in $b_F' = b_E + v > b_F$ which results in the step being rejected. \square

Corollary 3.5. *The RLS is stuck in a local optimum if $b_F - b_E < w_n$ holds and $b_F > opt$.*

Proof. A single bit flip of weight v can only happen if $b_F - b_E \geq v$. If $b_F - b_E < w_n$ there is no weight which satisfies the condition and therefore no single bit flip is possible. Since the RLS can only move one bit at a time and only if it results in an improvement, the RLS is stuck. \square

Corollary 3.6. *Every object $\leq \frac{W}{3}$ can be moved from b_F to b_E if $b_F \geq \frac{2W}{3}$*

Proof. $b_F \geq \frac{2W}{3} \Rightarrow b_E \leq W - \frac{2W}{3} \leq \frac{W}{3} \Rightarrow b_F - b_E \geq \frac{2W}{3} - \frac{W}{3} = \frac{W}{3}$ and with Lemma 3.4 the assumption follows. \square

Lemma 3.7. *In expected time $\mathcal{O}(n \log n)$ the weight of the fuller bin can be decreased to $\leq \frac{2W}{3}$ by the RLS and the (1+1) EA if every object besides the biggest in the fuller bin is at most $\frac{W}{3}$ and $w_1 \leq \frac{W}{2}$.*

Proof. In expected time $\mathcal{O}(n \log n)$ the RSH can move every object $\leq \frac{W}{3}$ to the emptier bin as long as $b_F \geq \frac{2W}{3}$ due to Corollary 3.6 and Theorem 3.1. So in expected time $\mathcal{O}(n \log n)$ the solution can be shifted to w_1 being in one bin and all other objects in the other bin. The RSH will only stop moving the elements if the condition $b_F \geq \frac{2W}{3}$ is no longer satisfied (Corollary 3.6). If $w_1 \geq \frac{W}{3}$ and every object was moved to the bin without w_1 , then $b_F = \max\{W - w_1, w_1\} = W - w_1 \leq \frac{2W}{3}$, because $w_1 \leq \frac{W}{2}$. So either the RSH moves all objects to the emptier bin or stops moving objects because $b_F < \frac{2W}{3}$ both resulting in $b_F \leq \frac{2W}{3}$. If w_1 is not in the fuller bin, then the result follows by Corollary 3.3.

Now assume $w_1 < \frac{W}{3}$. In this case the RLS will move one object per step to the emptier bin. Each object has weight $< \frac{W}{3}$ and therefore one step cannot decrease the weight of the fuller bin from $> \frac{2W}{3}$ to $\leq \frac{2W}{3}$. If all objects except one were moved to one bin, the other bin would have a weight of at least $W - w_1 > \frac{2W}{3}$. Therefore the RLS will find a solution with $b_F < \frac{2W}{3}$ before moving all elements from the first to the second bin.

The proof for the (1+1) EA is mostly the same. The main difference is the (1+1) EA being able to flip more than one bit in a single step. Such a step could make the emptier bin the fuller bin or increase the number of bits that must be shifted to the emptier bin. But with the results of Theorem 3.1 the proof works exactly the same as for the RLS. The case $w_1 \geq \frac{W}{3}$ does not change only the bin containing w_1 might change. Apart from that there is no difference for the (1+1) EA. The case $w_1 < \frac{W}{3}$ is also rather similar. The (1+1) EA will move elements from the fuller bin to the emptier bin until $b_F < \frac{2W}{3}$ holds. The (1+1) EA can make the emptier bin the fuller bin by moving multiple objects in one step, but this does not hinder it from reaching $b_F < \frac{2W}{3}$. After such a step it will continue moving elements until the condition holds. \square

Lemma 3.8. *The RLS and the (1+1) EA reach an approximation ratio of at most $\frac{4}{3}$ in expected time $\mathcal{O}(n \log n)$ if $w_1 < W/2$*

Proof. If $w_1 + w_2 > \frac{2W}{3}$ after time $\mathcal{O}(n)$ w_1 and w_2 are separated and will remain separated afterwards (3. Average case analysis, Theorem 1 in [Die05]). From then on the following holds. If w_1 is in the emptier bin, then the result follows directly by Corollary 3.3. Otherwise all elements in the fuller bin except w_1 have a weight of at most $\frac{1}{3}$ and therefore the result follows by Lemma 3.7 and Lemma 3.2. If $w_1 + w_2 \leq \frac{2W}{3}$ the result follows directly by Lemma 3.7 and Lemma 3.2. \square

Corollary 3.9. *The RLS and the (1+1) EA reach an approximation ratio of at most $\frac{4}{3}$ in expected time $\mathcal{O}(n \log n)$*

Proof. This follows directly from Theorem 3.1 and Lemma 3.8. \square

3.2. Binomial distributed input

Lemma 3.10. *A binomial distributed input $\sim B(m, p)$ has a perfect partition ($b_F - b_E = 0$ for even W and $b_F - b_E = 1$ for uneven W) with high probability if the input size n is large enough.*

Proof. Sketch:

- The initial distribution is likely rather close to the optimum
- The difference between the bins is probably not more than 10 expected values
- the large values

Consider a random separation of all values into two sets with equal size if n is even or one set with one value more than the other if n is odd. The sum X of one set is a sum of $\frac{n}{2} \cdot m$ independent Bernoulli trials with probability p . With Chernoff Bounds the following inequality follows:

$$\mathbb{P}(X \geq (\frac{n}{2} + \sqrt{\frac{n}{2}}) \cdot m \cdot p) = \mathbb{P}(X \geq (1 + 2\sqrt{\frac{2}{n}}) \cdot \frac{npm}{2}) \leq e^{-\frac{mnp}{2} \cdot 2\sqrt{\frac{2}{n}}/3} = e^{-\frac{2mp}{3}}$$

For $mp \geq 1.5$ the probability is less than $\frac{1}{e}$. Otherwise the input is rather trivial, since the numbers will be concentrated around 1.5 and most values will be below 10.

After moving $\mathcal{O}(\sqrt{\frac{n}{2}}/2)$ objects to the emptier set, the difference between the two sets is at most half the expected value mp of a single value. ... \square

Lemma 3.11. *With high probability the RLS does not find an optimal solution for an input with distribution $\sim B(m, p)$ if n and m are large enough.*

Proof. Sketch:

- There exists an optimal solution with high probability due to last lemma
- probability for a value to be very low is almost 0 if m is huge
- The RLS only moves one element per step and will reach $b_F - b_E < w_n$ without $b_F = opt$ being true
- \rightarrow RLS can't make another step and is stuck in a local optimum.

Due to Lemma 3.10 the input has an optimal solution with high probability. The probability for small values ... \square

4. Higher mutation rates and heavy tailed mutations

In this chapter different ways to change the mutation rate for EAs are discussed so that they flip more bits in expectation. For OneMax and all other linear functions the mutation rate of $p_m = 1/n$ was proven to be optimal[Wit13]. This is not the case for every fitness function. Jump_k has a optimal mutation rate of k/n and a small constant factor deviation from k/n results in an increase of the runtime exponential in $\Omega(k)$ [DLMN17].

4.1. Algorithms

For the (1+1) EA changing the expected amount of flipped bits per step can be done easily. By changing the mutation rate $1/n$ to c/n for any constant c the algorithm now flips $n \cdot c/n = c$ bits in expectation.

For the RLS it is not that simple, as the RLS chooses a random bit and flips it. Instead of flipping c bits in every step there should be the possibility to flip different amounts of bits in every step. The standard RLS chooses a random neighbour with Theorem 3.1 one. So the variant of the RLS could simply choose neighbours that have a Theorem 3.1 larger than one. The selection should still be uniform random to keep the idea of the RLS intact. One possible way is to choose a random neighbour with Theorem 3.1 $\leq k$. The amount of neighbours with Theorem 3.1 y is given by $\binom{n}{y}$. For $k = 4$, this results in n neighbours with Theorem 3.1 1, $n(n-1)/2$ neighbours with Theorem 3.1 2, $n(n-1)(n-2)/6$ and

Algorithm 4.1: (1+1) EA WITH STATIC MUTATION RATE

```
1 choose  $x$  uniform from  $\{0, 1\}^n$ 
2 while  $x$  not optimal do
3    $x' \leftarrow x$ 
4   flip every bit of  $x'$  with probability  $c/n$ 
5   if  $f(x') \leq f(x)$  then
6      $x \leftarrow x'$ 
```

Algorithm 4.2: RLS_k^B

```

1 choose x uniform from  $\{0, 1\}^n$ 
2 while x not optimal do
3    $x' \leftarrow$  uniform random neighbour of x with Theorem 3.1  $\leq k$ 
4   if  $f(x') \leq f(x)$  then
5      $x \leftarrow x'$ 

```

Algorithm 4.3: RLS_k^S

```

1 choose x uniform from  $\{0, 1\}^n$ 
2 while x not optimal do
3    $y \leftarrow$  uniform random value  $\in \{1, \dots, k\}$ 
4    $x' \leftarrow$  uniform random neighbour of x with Hamming Distance y
5   if  $f(x') \leq f(x)$  then
6      $x \leftarrow x'$ 

```

$n(n-1)(n-2)(n-3)/24$. The probability to choose a random neighbour with Theorem 3.1 $y \leq k$ for $k = \mathcal{O}(1)$ is given by

$$P(\text{RLS}_{k,k}^B \text{ flips } y \text{ bits}) = \frac{\binom{n}{y}}{\sum_{i=1}^k \binom{n}{i}} = \frac{\Theta(n^y)}{\sum_{i=1}^k \Theta(n^i)} = \frac{\Theta(n^y)}{\Theta(n^k)} = \Theta(n^{y-k}) = \Theta\left(\frac{1}{n^{k-y}}\right)$$

This variant of the RLS is likely to choose a neighbour with Theorem 3.1 k as the number of neighbours with hamming distance k rises with k for $k \leq n/2$. The probability of flipping only one bit is $\mathcal{O}(\frac{1}{n^{k-1}})$. For some inputs flipping only one bit might be more optimal which is rather unlikely for this variant of the RLS which will be called RLS_k^B from now on.

An alternative way of changing the RLS is to first choose $y \in \{1, \dots, k\}$ uniform random and then choose a neighbour with Theorem 3.1 y uniform random. Here the probability of flipping $y \leq k$ bits is given by $1/k$, so the algorithm is much more likely to choose to flip only one bit. This variant of the RLS will be referred to as RLS_k^S .

Both variants of the RLS change at most k bits in each step and therefore only a constant amount of bits. For the (1+1) EA the algorithm will also flip mostly $\mathcal{O}(c)$ bits which is also constant. So neither of the new variants is likely to change up to $\mathcal{O}(n)$ bits. Quinzan *et al.* therefore introduced another mutation operator in [FGQW18] called pmut_β . This operator chooses k from a powerlaw distribution D_n^β with exponent β and maximum value n and then k uniform random bits are flipped. This algorithm will mostly flip a small number of bits but occasionally up to n bits. Distributions like this are called heavy tailed mutations because their tail is not bounded exponentially.

4.2. Runtime Analysis of higher mutation rates

The first analysed input is again the input with $w_1 \geq W$. This input is very close to a weighted OneMax and therefore more or less the OneMax equivalent for Partition. The first analysed algorithm will be the RLS variants.

Lemma 4.1. *The expected time until the RLS_k^S and RLS_k^B flips w_1 after they reached a solution with $b_E = c \cdot \frac{W-w_1}{2}$ with $1 < c < 2$ is at least $\frac{n(n-1)(c-1)}{k(k-1)^2}$. For constant values of c the expected time is $= \Omega(n^2)$.*

Proof. For a successful flip of w_1 after $b_E \geq \frac{W-w_1}{2}$ holds a total volume of $2 \cdot (b_E - \frac{W-w_1}{2})$ must be shifted from b_E to b_F . Otherwise the step is rejected. If the algorithm moves $2 \leq y \leq k$ elements in this step there are $y-1$ bits left to shift the volume from b_E to b_F . The object with the highest shifted volume must have a volume of at least $2 \cdot (b_E - \frac{W-w_1}{2}) / (y-1)$ because otherwise there is not enough total volume shifted from b_E to b_F . For $b_E = c \cdot \frac{W-w_1}{2}$ with $1 < c < 2$ this leads to at most $d \leq \frac{W-w_1}{2(c \cdot \frac{W-w_1}{2} - \frac{W-w_1}{2}) / (y-1)} = \frac{(W-w_1)(y-1)}{(W-w_1)(c-1)} = \frac{y-1}{c-1}$ objects with the given weight. The probability for such a step, given the algorithm flips y bits, is therefore at most

$$\frac{\binom{1}{1} \cdot \binom{d}{1} \cdot \binom{n-2}{y-2}}{\binom{n}{y}} = \frac{d \frac{(n-2)!}{(n-2-y+2)! \cdot (y-2)!}}{\frac{n!}{(n-y)! \cdot y!}} = \frac{d \cdot (n-2)! \cdot (n-y)! \cdot y!}{n! \cdot (n-y)! \cdot (y-2)!} = \frac{dy(y-1)}{n(n-1)}$$

The expected time for such a step to happen is at least

$$\frac{1}{\mathbb{P}(y \text{ bits are flipped})} \cdot \frac{n(n-1)}{dy(y-1)} \geq \frac{1}{1} \cdot \frac{n(n-1)(c-1)}{y(y-1)^2} \geq \frac{n(n-1)(c-1)}{k(k-1)^2}$$

k is constant and if c is constant too, this leads to expected time $= \Omega(n^2)$ for a flip of the first bit if $b_E = c \cdot \frac{W-w_1}{2}$ with $1 < c < 2$. \square

Lemma 4.2. *Let $w_1 \geq W/2$, then for any $\gamma > 1$ and $0 < \delta < 1$, the $(1+1)$ EA (RLS_k^S) reaches an f -value at most $w_1 + \delta(W-w_1)$ in at most $\lceil en \ln(\gamma/\delta) \rceil$ ($\lceil kn \ln(\gamma/\delta) \rceil$) steps with probability at least $1 - \gamma^{-1}$. Moreover, the expected number of steps is at most $2\lceil en \ln(2/\delta) \rceil$ ($2\lceil kn \ln(2/\delta) \rceil$).*

Proof. This Lemma is almost the exact same as C. Witts Lemma 2 from section 2. Definitions and Proof Methods in [Die05]. The proof is almost the same. Instead of choosing the lower bound $W/2$ as p_0 here the exact value $W - w_1 \leq W/2$ is chosen. For the RLS_k^S setting $t' = kn \ln(\gamma/\delta)$ suffices to show the result. \square

Lemma 4.3. *The $RLS_{k=2}^S$ reaches the optimal solution on an input with $w_1 \geq W/2$ in expected time $\mathcal{O}(n \log n)$*

Proof. The run can be divided in the same two phases as in Theorem 3.1. The expected length of the first phase is at most kn in expectation because the probability of flipping only the first bit is $1/kn$. Flipping the first bit together with other bits might be successful as well, but it will not be accepted in every case.

In the second Phase the RLS_k^S tries to remove every element of the bin containing w_1 . In contrast to the RLS it can make steps increasing the hamming distance to the optimum once the algorithm is in phase 2. Steps where the first bit is not flipped can increase the Hamming distance at most by $k-1$, since the global optimum cannot change without changing w_1 and moving all k elements in the wrong direction will always be rejected. After expected time $(2\lceil kn \ln(2/0.4) \rceil) = 2\lceil kn \ln(5) \rceil \leq 2\lceil 1.61 \cdot kn \rceil \leq 4kn$ the RLS_k^S reaches a solution of $f(x) \leq w_1 + 0.4(W-w_1)$ (Lemma 4.2). This means $b_E \geq 0.6(W-w_1) = 1.2 \cdot \frac{W-w_1}{2}$. With Lemma 4.1 the expected time of at least $\frac{n(n-1)(c-1)}{k(k-1)^2} = \Omega(n^2)$ for a successful flip of w_1 follows.

The RLS_2^S chooses to flip only one bit with probability $1/2$. In these steps it behaves exactly the same as the standard RLS. For $k=2$ a step moving two elements cannot increase the Hamming distance if w_1 is not flipped. Using a fitness level argument the expected running time of phase 2 if steps flipping the first bit are always rejected is at most

$$E(T_{Phase2}) \leq \sum_{i=1}^n \frac{kn}{i} = kn \cdot \sum_{i=1}^n \frac{1}{i} \leq kn(\ln(n) + 1) = kn \ln(n) + kn = \mathcal{O}(n \ln(n))$$

So the algorithm reaches an optimal solution in expected time $\mathcal{O}(n \ln(n))$ whereas a second flip of w_1 only happens in expected time $\Omega(n^2)$. The algorithm therefore reaches an optimal solution in expected time

$$E(T_{Phase1}) + E(T_{Phase2}) \leq 4kn + kn \ln(n) + kn = kn \ln(n) + 5kn = \mathcal{O}(n \ln(n))$$

□

5. Experimental Results

In the following chapter the different variants of the RLS and the (1+1) EA are now analysed empirically for the best algorithm depending on the input. Additionally for most lemmas from the previous chapters there are also tests if they actually hold in practice.

5.1. Code

The complete java code used for all empirical studies is available on GitHub.

5.1.1. The Algorithms

All different variants of the RLS function more or the less the same. They start with an initial random value and then optimise this one value in the loop. The loop can be summarised like this:

1. generate a number k of bits to be flipped (algorithm specific)
2. flip k random bits
3. evaluate fitness of the mutated individual
4. replace old value with new value if new value is better
5. repeat if not optimal

The (1+1) EA variants behave differently at first glance as they flip bit each bit independently with probability c/n . This can be seen as n independent Bernoulli trials with probability c/n . The amount of bits that are flipped is therefore binomial distributed and the algorithm can be implemented exactly as the versions of the RLS. The same holds for the *pmut* operator which generates a number k from a powerlaw distribution and then flips k bits. This leads to only one implementation of a partition solving algorithm which is not only given the input array of numbers but also a generator for the amount of bits to be flipped in each step. The random values for the amount of bits to be flipped are generated according to this table:

| Algorithm | Returned value |
|-------------|---|
| RLS | 1 |
| RLS_k^B | $y \in \{1, \dots, k\}$ with probability $\frac{\binom{n}{y}}{\sum_{i=1}^k \binom{n}{i}}$ |
| RLS_k^S | uniform random value $y \in \{1, \dots, k\}$ |
| (1+1) EA | binomial distributed value from $\sim B(n, c/n)$ |
| <i>pmut</i> | powerlaw distributed value from $\sim D_n^\beta$ |

Algorithm 5.1: GENERICPARTITIONSolver

```

1 choose x uniform random from  $\{0, 1\}^n$ 
2 while x not optimal do
3    $x' \leftarrow x$ 
4    $k \leftarrow \text{kGenerator.generate}()$ 
5   flip k uniform random bits of  $x'$ 
6   if  $f(x') \leq f(x)$  then
7      $x \leftarrow x'$ 

```

5.1.2. Random number generation

Java only provides a random number generator for uniform distributed values for any integer interval or random double values $\in [0, 1)$. For this project this does not suffice as for an efficient way of implementing the (1+1) EA or simply for generating a binomial distributed input another random number generator is needed. One of the needed distributions is a binomial distribution. The simplest way to generate a number $\sim B(m, p)$ would be to run a loop m times and add 1 to the generated number if a uniform random value $\in [0, 1)$ is less than p . This works perfectly fine and generates numbers according to the distribution. With low values for p this approach is rather inefficient and especially for values of $p = 1/m$. The expected value in this case is 1 but generating a random number takes time $\mathcal{O}(m)$. Another more efficient way was implemented by StackOverflow user pjs on stackoverflow inspired by Devroyes method introduced in [Dev06]. This method has an expected running time of $\mathcal{O}(mp)$ which is equal to the expected value of the distribution. For the case of $p = 1/m$ this runs in expected constant time in comparison to $\mathcal{O}(m)$ for the naive way. This number generation was also used for the implementation of the (1+1) EA instead of running a loop in every step.

TODO: explain how and why this algorithm works

Algorithm 5.2: BINOMIAL RANDOM NUMBER GENERATOR

```

1  $q \leftarrow \ln(1.0 - p)$ 
2  $x \leftarrow 0$ 
3  $sum \leftarrow 0$ 
4 while true do
5    $sum \leftarrow sum + \ln(\text{random}()) / (n - x)$ 
   // random() generates a random value  $\in [0, 1)$ 
6   if  $sum < q$  then
7     return x
8    $x \leftarrow x + 1$ 

```

The next generator needed is for geometric distributed values. This generator is only necessary for the generation of geometric distributed inputs but not for the algorithms. The easiest way to generate geometric distributed values is the naive way: generating a uniform random value p' until $p' < p$ holds. The expected running time of this algorithm is equal to the expected value of the distribution $1/p$. So this method is comparably effective to the approach used for binomial random number generation.

The last generator needed is for powerlaw distributed values. This generator is in contrast to the geometric number generator needed for both the algorithm with the $pmut_\beta$ mutation

Algorithm 5.3: GEOMETRIC RANDOM NUMBER GENERATOR

```

1  $sum \leftarrow 0$ 
  // random() generates a random value  $\in [0, 1)$ 
2 while  $random() \geq q$  do
3    $sum \leftarrow sum + 1$ 
4 return  $sum$ 

```

operator and for generating inputs. This implementation is also from stackoverflow. The user gnovice provided the following formula on this page on stackoverflow:

$$x = [(b^{n+1} - q^{n+1}) * y + a^{n+1}]^{1/(n+1)}$$

a is the lower bound, b the upper bound, n the parameter of the distribution and y the number generated uniform random $\in [0, 1)$. The idea behind the formula and the formula itself is explained in a mathworld page.


TODO: give a brief explanation here

5.2. Do inputs have perfect partitions?

5.2.1. Binomial Inputs

Lemma 3.10 is only valid for larger n . In practice the bound is much smaller depending on the expected value of a single number. Another factor deciding how likely an input is to have a perfect partition is whether n is even or odd. To determine the influence of all factors multiple experiments were conducted. The goal of the first experiment was to determine the influence of the array size to the input having a perfect partition and the fact if n is even or odd. So for every possible combination of $p \in \{0.1, 0.2, \dots, 0.8, 0.9\}$, $m \in \{10, 100, 1000, 10^4, 10^5\}$ and $n \in \{2, 3, 4, \dots, 19, 20\}$ 1000 randomly generated inputs of size n were tested for a perfect partition. Due to the small values for n it was possible to brute force the results in a short amount of time. The results are visualised in figure 5.1 to figure 5.9.

Figure 5.1.: Percentage of Binomial inputs with perfect partitions for $m = 10$



On the x -axis is the size of the input and on the y -axis the percentage of inputs that had a perfect partition. The different graphs in each figure resemble the different values of p used for generating the inputs. The graph for $p = 0.1$ resembles the percentage of inputs that had a perfect partition with values generated from the distribution $\sim B(m, 0.1)$ with m being dependent on the figure. For figure 5.1 m has the value 10.

Figure 5.1 is a bit overloaded with information and the zigzag makes it hard to gain any benefit from the graphs. That's why for $m \in \{10, 100, 1000\}$ there is one figure for the even input sizes n and one for the odd. For $n \geq 10,000$ all values for the odd input sizes are 0

%, so there is no point in showing the data in a separate graph. The lines in the separate graphs for even and odd n are only to visualise the growth in n . The displayed values in between the listed input sizes on the x -axis do not have a meaning.

TODO: rephrase this sentence to make it more obvious what is meant

It is easy to see that for small inputs sizes it is relevant if n is even or odd for higher expected values as all curves in figure 5.9 oscillate between 0% and 100% for $n \geq 14$. For odd inputs the probability of a perfect partition decreases much more drastically with m as for even inputs because the expected value of a single number increases with m . If all values are much higher the small differences between the values can no longer even out the fact of one set having more elements than the other. The oscillation therefore increases with increasing m . For $n = 20$ all 1000 inputs had a perfect partition for every combination of p and m but for $n = 19$ only combinations where $mp \leq 300$ holds had at least one input with a perfect partition. For expected values of up to 10^5 it seems to be almost granted that an input of length 20 has a perfect partition if it is binomial distributed. Even for only 12 binomial generated values more than 50% of the inputs had a perfect partition (see figure 5.9). Another visible effect is the decreasing percentage with rising p . This may be a direct result of the value chosen for p but can also be an indirect result as the value for p changes the expected value for a constant m . The expected value may have an influence on the number of perfect partitions because it influences the highest value of the input. For uniform distributed inputs Borgs *et. al.* showed that the coefficient of number of bits needed to encode the max value/ n has a huge impact on the number of perfect partitions [BCP01]. For a coefficient < 1 the probability of a perfect partition tends to 1 and for a coefficient > 1 it tends to 0. This was only proven for the uniform distributed input, but it might also hold for a binomial distributed input. This leads to the second experiment.

Figure 5.2.: Percentage of Binomial inputs with perfect partitions for $m = 10$ for even n




Figure 5.4.: Percentage of Binomial inputs with perfect partitions for $m = 100$ for even n




Figure 5.3.: Percentage of Binomial inputs with perfect partitions for $m = 10$ for odd n









Figure 5.5.: Percentage of Binomial inputs with perfect partitions for $m = 100$ for odd n



In the second experiment the inputs were generated a bit differently. Here the goal was to keep the expected value fixed for any combination of p and n and set the value of m

Figure 5.6.: Percentage of Binomial inputs with perfect partitions for $m = 1000$ for even n

 Figure 5.7.: Percentage of Binomial inputs with perfect partitions for $m = 1000$ for odd n

 Figure 5.8.: Percentage of Binomial inputs with perfect partitions for $m = 10,000$

 Figure 5.9.: Percentage of Binomial inputs with perfect partitions for $m = 100,000$


to e/p for all $e \in \{10, 20, 30, 40, 50, 100, 200, 500, 1000, 2000, 5000, 10000, 50000\}$ so that $E(X) = mp = e/p \cdot p = e$. With this setup the influence of the expected value is almost isolated from the other parameters. The probability is still linked to p as p also influences the variance $mp(1-p)$.

Figure 5.10 to figure 5.13 again show the percentage of perfect inputs with different settings of m, p, n . The x -axis is the expected value mp of a single number of the input. The different graphs show the percentage for different input sizes. It seems as if the value of p has a much smaller influence than the expected value. For a fixed expected value and a fixed input size a higher value for p seems to only slightly increase the percentage of inputs with a perfect partition. The expected value influences the percentage significantly more. For $p = 0.1, n = 14$ the value decreases from 100% at $E(X) = 10$ to below 20% at $E(X) = 50000$ (figure 5.10). For $p = 0.9$ the percentage only drops below 50% but still decreases by a factor of 2 (figure 5.13).

The last experiment showed that for $n = 20$, 1000/1000 inputs had a perfect partition. This raised the question of how the amount of perfect partition changes with changing values for m, p, n . Figure 5.14 to figure 5.17 show the amount of perfect partitions a binomial distribution $\sim B(m, p)$ has. For these figures 10,000 random binomial inputs with the given values for m and p were generated. Each input was then tested for the number of perfect partitions it has. The used method was again brute force to ensure correctness which was only possible due to the small input sizes. After all runs the average values were combined in the given figures. The value of p is dependent on the picture and each value of $m \in \{10, 100, 1000, 10000\}$ has its own graph within the figure. The x -axis is the size of the input and the y -axis the number of perfect partitions the input has. Notice that all graphs have a y -axis with a logarithmic scale. Since the graphs are all linear the actual values rise exponentially. The number of perfect partitions is mostly multiplied by a factor between 3 and 4 when the input size increases by 2.

Figure 5.10.: Percentage of Binomial inputs with perfect partitions for $p = 0.1$




Figure 5.11.: Percentage of Binomial inputs with perfect partitions for $p = 0.2$




Figure 5.12.: Percentage of Binomial inputs with perfect partitions for $p = 0.5$





Figure 5.13.: Percentage of Binomial inputs with perfect partitions for $p = 0.9$







The higher the value of m the closer the curves of p and $1 - p$ get. For $m = 1000$ and $m = 10000$ the values are almost the same for every input size. For $p = 0.1, m = 10$ an input with expected values of 1 seems to be much more likely to have a perfect partition than an input with expected value $10 \cdot 0.9 = 9$. With growing m this has less impact.

So the binomial input should be easy to solve due to the exponential number of perfect partitions. It might be harder for the smaller values of n as there are only a few perfect partitions. Due to the small number of total possibilities it should still be easy to solve for the small values of n as long as the RSH is not stuck in a local optimum. The number of iterations might be high in terms of the big-O notation but should still be small in the absolute value.

5.3. Binomial distributed values

In the following subsections the performance of the different algorithms is tested for different kinds of inputs. The exact distributions of the input are explained separately in each subsection. The procedure for each comparison is always the same. A random input is generated according to the distribution and then solved by every algorithm. All algorithms had the same two stopping conditions. The first was reaching a perfect partition and the second was taking more than $10 \cdot n \ln(n)$ steps or $100 \cdot n \ln(n)$ in some cases. For the lower values of n the step limit of 100,000 was used instead. For $n = 20$ giving the algorithm

Figure 5.14.: Amount of perfect partitions for $p = 0.1$ Figure 5.15.: Amount of perfect partitions for $p = 0.9$ Figure 5.16.: Amount of perfect partitions for $p = 0.2$ Figure 5.17.: Amount of perfect partitions for $p = 0.5$ 

only 600 steps is rather small. In some cases the smaller inputs are even more difficult to solve. Most modern computer should be able to handle 100,000 iterations in a short amount of time anyway. So the minimum step limit of 100,000 seemed reasonable. If either of these conditions was met, the algorithm returned its current best solution. This step is repeated 1000 times. The results are presented in a table containing multiple statistics for each algorithm over all 1000 runs. The data is explained in the table below.


| column name | meaning |
|-----------------|--|
| algo type | type of algorithm (RLS, RLS_k^B , RLS_k^S , (1+1)EA or pmut) |
| algo param | parameter of the algorithm or '-' if it is the standard variant |
| avg mut/change | average #bits flipped for iterations leading to an improvement |
| avg mut/step | average #bits flipped for any iteration |
| total avg count | average #iterations for all runs |
| avg eval count | average #iterations of runs returning an optimal solution |
| max eval count | maximum #iterations of runs returning an optimal solution |
| min eval count | minimum #iterations of runs returning an optimal solution |
| fails | number of runs that did not find an optimal solution |
| fail ratio | ratio of unsuccessful runs to all runs |
| avg fail dif | average value of $b_F - f(opt)$ for non-optimal solutions |

Firstly the different variants of the RLS are compared with values of $k \in \{2, 3, 4\}$, then the performance of the (1+1) EA with static mutation rate c/n with $c \in \{1, 2, 3, 5, 10, 50, 100\}$ and lastly the performance of the $pmut_\beta$ mutation operator with the parameter $\beta \in \{1.25, 1.5, \dots, 2.75, 3.0, 3.25\}$. Additionally the best variants of each algorithm are compared in another 1000 runs. Afterwards there is also a comparison for multiple input sizes of the best algorithms because the best algorithm is often dependent on the size of the input. Normally there are three tables for each input. The first states how often the algorithms did not find an optimal solution for the different input sizes ('fails' in top left cell). The

second gives their average performance for the successful runs ('avg' in top left cell) and the last the performance for all runs ('total avg' in top left cell). The last two tables differ in the unsuccessful runs. Often the algorithm is stuck in a local optima it won't leave in reasonable time or never for variants of the RLS. In these cases the step limit is the deciding factor on how big the penalty for this run is. So neither of the two average values alone is enough to give a complete insight on the performance. Sometimes a variant of the RLS is much faster than the other algorithms for a specific input but is also the only algorithm to get stuck in a local optimum. This creates the possibility to start the RLS variant with a low step limit and switch to the (1+1) EA if the RLS variant does not return an optimal solution. Giving both tables for the different average values might help with this decision.

The first analysed inputs are inputs following a binomial distribution $\sim B(m, p)$ as those inputs have been researched in the previous subsection. The results showed that the expected value of a single number is the main driver for the amount of perfect partitions the input has. The results also suggested the inputs tend to have more perfect partitions if the expected value is lower. The more perfect partitions an input has relative to the number of all possible partitions, the more likely the different RSHs are to find one of those. Therefore researching inputs with higher expected values seems more interesting but generating higher values takes more time with a random number generator that needs $\mathcal{O}(mp)$ time. To keep the time for generating one set of numbers reasonable the values chosen for all tests are $m = 10000, p = 0.1, n = 10000$ with the expected value for a single number being $mp = 1000$. Figure 5.18 shows a random binomial distributed input of length $n = 10000$. For this input type almost every time all elements were sharply concentrated around the expected value with all values being at 1000 ± 200 (for figure 5.18 even closer at 1000 ± 121). So after reaching a difference between the two bins of below $(1000 - 200)/2 = 400$ the algorithm can no longer achieve an improvement by flipping a single bit.

Figure 5.18.: Distribution of a random binomial input



5.3.1. RLS Comparison

| algo type | RLS_k^B | RLS_k^B | RLS_k^S | RLS_k^S | RLS_k^S | RLS_k^B | RLS |
|-----------------|-----------|-----------|-----------|-----------|-----------|-----------|---------|
| algo param | b=2 | b=4 | s=2 | s=4 | s=3 | b=3 | - |
| avg mut/change | 2.000 | 4.000 | 1.603 | 2.553 | 2.000 | 2.728 | 1.000 |
| avg mut/step | 2.000 | 4.000 | 1.500 | 2.500 | 1.999 | 3.000 | 1.000 |
| total avg count | 318 | 434 | 499 | 579 | 681 | 518,428 | 920,109 |
| avg eval count | 318 | 434 | 499 | 579 | 681 | 395,440 | 50 |
| max eval count | 1,648 | 3,243 | 3,094 | 3,737 | 4,717 | 917,134 | 50 |
| min eval count | 20 | 28 | 11 | 17 | 16 | 0 | 50 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.234 | 0.999 |
| avg fail dif | - | - | - | - | - | 1 | 254 |

The RLS_2^B seems to perform the best as it mostly switches two elements which works great for binomial distributed inputs. The same algorithm with $k = 4$ performs a bit worse but still good as switching 4 elements can be beneficial as well. The variant of RLS_3^B on the other hand does not reach the optimal solution in 23.4% of the inputs with an average difference of 1. It also needs 1000 times more iterations to find an optimum on average compared to the best algorithm RLS_2^B . The RLS_k^S variants behave mostly the same with $k = 2$ being the best, followed by $k = 4$ and $k = 3$. In this case the variant of $k = 3$ is by far not as bad as for the RLS_3^B because the probability of flipping 2 bits is $1/3$ as compared to $\mathcal{O}(n^{-1})$ for the RLS_3^S . The RLS_k^S seem all to be a good option for binomial inputs with values of $k \in \{2, 3, 4\}$. The standard RLS on the other hand performs by far the worst as it only moves one element per step. It only managed to reach the optimal solution once for 1000 different inputs. The number of iterations for this input was only 50 so the RLS likely had a good initialisation with a few lucky steps leading directly to the optimum. For all other cases the average difference between the bins was 254 which is close to the median of the values from 0 to $(1000 - 100)/2 = 450$. This is likely due to the RLS being unable to improve the solution once the current solution has a difference below half of the lowest value (Corollary 3.5).

5.3.2. (1+1) EA Comparison

For the (1+1) EA the best static mutation rate seems to be $3/n$. The probability of flipping 2 or 4 bits as n goes to infinity for mutation rate $1/n$ approaches $13/24e \approx 0.199$, for $2/n$ approaches $8/3e^2 \approx 0.361$, for $3/n$ approaches $63/8e^3 \approx 0.392$, for $4/n$ approaches $56/3e^4 \approx 0.342$ and for $5/n$ approaches $77/2e^5 \approx 0.259$. So the highest probability has $c = 3$, followed by $c = 4$ and $c = 2$ then $c = 5$ and lastly $c = 1$. For higher values of c the probability decreases further as the expected number of flipped bits is c for mutation rate c/n .

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
|----------------|-------|-------|-------|-------|-------|--------|--------|---------|
| algo param | $3/n$ | $4/n$ | $2/n$ | $5/n$ | - | $10/n$ | $50/n$ | $100/n$ |
| avg mut/change | 3.101 | 3.968 | 2.343 | 4.859 | 1.698 | 9.732 | 49.544 | 99.494 |
| avg mut/step | 2.999 | 4.003 | 2.002 | 4.999 | 1.001 | 9.998 | 49.998 | 99.997 |
| avg eval count | 646 | 701 | 706 | 857 | 1,123 | 1,508 | 8,175 | 15,485 |
| max eval count | 5,346 | 5,692 | 3,415 | 5,572 | 7,001 | 12,112 | 52,831 | 145,269 |
| min eval count | 23 | 4 | 30 | 9 | 23 | 14 | 27 | 69 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The static mutation rate $3/n$ seems to perform the best with both $4/n$ and $2/n$ being a close second place. The next best values are $5/n$ and the standard $1/n$ both having a clear difference between each other and the better parameters. From then on the number of iterations rises monotonically with rising mutation rate. The higher mutation rates perform significantly worse but the still find a solution within the limit as opposed to the standard RLS.

5.3.3. pmut Comparison

| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
|----------------|-------|-------|--------|-------|-------|-------|-------|--------|---------|
| algo param | 2.25 | 2.00 | 1.75 | 2.50 | 2.75 | 3.00 | 3.25 | 1.50 | 1.25 |
| avg mut/change | 3.822 | 6.266 | 14.371 | 2.804 | 2.347 | 1.995 | 1.843 | 38.318 | 92.365 |
| avg mut/step | 4.344 | 8.504 | 22.176 | 2.878 | 2.272 | 1.933 | 1.732 | 70.476 | 224.535 |
| avg eval count | 652 | 668 | 675 | 688 | 697 | 718 | 758 | 785 | 1,050 |
| max eval count | 4,340 | 4,506 | 5,616 | 5,098 | 9,140 | 5,081 | 6,189 | 6,542 | 7,837 |
| min eval count | 14 | 4 | 9 | 12 | 27 | 10 | 11 | 21 | 7 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

For the pmut_β mutation operator the choice of β seems to be much more insignificant than for the RLS or (1+1) EA. Here all values perform comparably good with only the value of

$\beta = 1.25$ having a clear performance difference compared to next best value. All values of β reach an optimal solution in every case. The worst variant of the $pmut_\beta$ operator still performs much better than the worst value for the (1+1) EA and even better than the worst RLS variant. There is no clear winner but because $\beta = 2.25$ had the best performance in this experiment, it was used for the comparison of the best variants.

5.3.4. Comparison of the best variants

| | | | |
|----------------|-------------------------------|-------|-------|
| algo type | RLS _k ^B | EA-SM | pmut |
| algo param | b=2 | 3/n | 2.25 |
| avg mut/change | 2.000 | 3.092 | 3.965 |
| avg mut/step | 2.000 | 2.999 | 4.339 |
| avg eval count | 302 | 677 | 691 |
| max eval count | 1,610 | 6,404 | 5,205 |
| min eval count | 9 | 33 | 17 |
| fail ratio | 0.000 | 0.000 | 0.000 |

TODO: ask what turning the tables into arrays means For this setting of $m = 10000, p = 0.1, n = 10000$ the RLS₂^B performs better than the (1+1) EA and $pmut_\beta$ mutation for all values of c/n and β by a factor of at least 2. This is likely from the fact that this version of the RLS flips almost only two bits which seems to be close to optimal for this kind of input. There are many values close to the expected value which can be switched to make small adjustments to the fitness value. The (1+1) EA with $p_m = 3/n$ and $pmut_\beta$ algorithm with $\beta = 2.25$ perform almost the same. To further investigate which input performs best on all binomial distributed inputs now a comparison with different input lengths follows. The parameters of the distribution were not changed.

The following table shows the number of runs in which the algorithms did not find an optimal solution within 50,000 steps. The time limit was set 50,000 because the algorithms normally reach the optimal solution within a few thousand steps. If the solutions is not found after 50,000 steps, the algorithm is most likely stuck in a local optimum which could only be left by flipping more bits than possible for the algorithm.

TODO: Change multipleN fails to show percentages instead of count

TODO: Add captions to multipleN tables

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|-----|----|-----|-----|------|------|-------|
| RLS ₂ ^S | 231 | 22 | 3 | 0 | 0 | 0 | 0 |
| RLS ₄ ^B | 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| RLS ₂ ^S | 243 | 4 | 0 | 0 | 0 | 0 | 0 |
| RLS ₄ ^S | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA (3/n) | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| pmut(2.5) | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

The (1+1) EA and $pmut_\beta$ always reach an optimal solution but the RLS does not. The RLS variants that can only flip two bits per step perform significantly worse for small inputs. They are probably more likely to get stuck in a local optimum where a step flipping 4 bits or more would be necessary. So the RLS₂^B does perform better for larger inputs but is much more likely to get stuck in a local optima. The next table contains the average number of iterations the algorithm needed to find an optimal solution for all runs where the algorithms managed to find an optimal solution. Here it still looks like the RLS₂^B finds the solution with the lowest amount of steps, because the cases where the algorithm is stuck in a local optima are not contained in this table.

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|------|-----|-----|-----|------|------|-------|
| RLS ₂ ^S | 164 | 257 | 337 | 256 | 257 | 293 | 310 |
| RLS ₄ ^B | 349 | 355 | 661 | 412 | 412 | 436 | 425 |
| RLS ₂ ^S | 267 | 435 | 408 | 428 | 442 | 452 | 496 |
| RLS ₄ ^S | 601 | 491 | 492 | 491 | 534 | 562 | 560 |
| (1+1) EA (3/n) | 716 | 570 | 569 | 580 | 614 | 625 | 650 |
| pmut(2.5) | 1596 | 576 | 600 | 637 | 657 | 700 | 685 |

The next table contains the overall average amount of iterations for every run. So runs where no optimal result was found add 50,000 to the sum of all iterations.


| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|-------|------|-----|-----|------|------|-------|
| RLS ₂ ^S | 11676 | 1351 | 486 | 256 | 257 | 293 | 310 |
| RLS ₄ ^B | 349 | 405 | 760 | 412 | 412 | 436 | 425 |
| RLS ₂ ^S | 12352 | 633 | 408 | 428 | 442 | 452 | 496 |
| RLS ₄ ^S | 650 | 491 | 492 | 491 | 534 | 562 | 560 |
| (1+1) EA (3/n) | 716 | 570 | 569 | 580 | 614 | 625 | 650 |
| pmut(2.5) | 1596 | 576 | 600 | 637 | 657 | 700 | 685 |

Here the RLS₂^B is only the best algorithm for values of $n \geq 500$. Below this bound choosing the (1+1) EA with static mutation rate $3/n$ is a safer choice as the (1+1) EA reaches an optimal solution for every input (in this experiment).

5.4. Geometric distributed values

For the geometric distribution the chosen default value is $p = 0.001$. This results in an expected value of 1000 which is the same as for the binomial distribution in the last subsection. This should make the results more comparable. The maximum value is theoretically not limited but for the implementation in Java the maximum value was set to the maximum value of a long value $= 2^{63} - 1 = 9,223,372,036,854,775,807$. Without this maximum value the value might overflow and instead be negative with high absolute value. Figure 5.19 shows a random geometric distributed input. The span of all values is way higher than for the binomial distribution, although they have same expected value. Here the values are not in the interval $[800, 1200]$ but rather between 0 and the 9000. The theoretical limitation of the values being at most $2^{63} - 1$ seems to not have an influence on the results.

Figure 5.19.: Distribution of a random geometric input



The geometric distribution does not only have low values close or equal to 1 but also has mostly values that are very small. This should lead to 1-bit flips being effective as the

small values can remove the small differences. Because there are so many small values moving only one bit might be better than switching two elements.

5.4.1. RLS Comparison

| algo type | RLS _k ^S | RLS _k ^S | RLS _k ^S | RLS _k ^B | RLS _k ^B | RLS _k ^B | RLS |
|-----------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|--------|
| algo param | s=2 | s=3 | s=4 | b=2 | b=3 | b=4 | - |
| avg mut/change | 1.477 | 1.959 | 2.431 | 2.000 | 3.000 | 4.000 | 1.000 |
| avg mut/step | 1.500 | 2.001 | 2.501 | 2.000 | 3.000 | 4.000 | 1.000 |
| total avg count | 2,592 | 2,945 | 3,259 | 3,497 | 4,463 | 5,345 | 6,650 |
| avg eval count | 2,592 | 2,945 | 3,259 | 3,497 | 4,463 | 5,345 | 2,055 |
| max eval count | 19,845 | 23,932 | 28,532 | 23,824 | 30,881 | 41,600 | 25,889 |
| min eval count | 8 | 22 | 19 | 18 | 43 | 19 | 23 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.005 |
| avg fail dif | - | - | - | - | - | - | 1 |

For these inputs the variants of the RLS perform differently to the binomial input. The only similarity is the RLS being the worst as the RLS is the only algorithm that did not find an optimal solution for every input. If the RLS did find an optimal solution in those 5 cases it would instead be the best RLS variant. The other algorithms are ranked by the probability of flipping only one bit. This means at first the three RLS_k^S variants from 2 to 3 to 4 and then the same for the RLS_k^B variants. So it does seem like moving mostly one element at once is better for the geometric input in comparison to two elements for the binomial distribution. In the 5 cases where the RLS did not find an optimal solution it was most likely stuck in a local optimum.

5.4.2. (1+1) EA Comparison

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
|----------------|--------|--------|--------|--------|--------|--------|---------|---------|
| algo param | 2/n | - | 3/n | 4/n | 5/n | 10/n | 50/n | 100/n |
| avg mut/change | 2.255 | 1.554 | 3.038 | 3.948 | 4.883 | 9.821 | 49.798 | 99.814 |
| avg mut/step | 2.000 | 1.000 | 3.000 | 4.001 | 5.000 | 9.999 | 49.998 | 100.001 |
| avg eval count | 3,712 | 3,833 | 4,195 | 4,472 | 5,465 | 8,282 | 21,648 | 29,404 |
| max eval count | 39,593 | 53,450 | 33,598 | 42,449 | 55,717 | 65,522 | 149,048 | 281,857 |
| min eval count | 18 | 13 | 15 | 14 | 25 | 23 | 46 | 17 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The results for the (1+1) EA are similar to the results of the RLS. From mutation rate 2/n on the runtime increases with rising mutation rate. The only part that does not fit into the theory of 1 bit flips being superior is the mutation rate 2/n performing better than the standard 1/n. The average number of iterations for the standard (1+1) EA is only slightly higher than for the mutation rate 2/n, so this might be just due to a too small number of runs of the algorithms. All variants reach an optimal solution within the given limit for the number of iterations.

5.4.3. pmut Comparison

| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
|----------------|--------|--------|--------|--------|--------|--------|--------|--------|---------|
| algo param | 3.25 | 3.00 | 2.50 | 2.75 | 2.25 | 2.00 | 1.75 | 1.50 | 1.25 |
| avg mut/change | 1.682 | 1.872 | 2.688 | 2.150 | 3.704 | 6.938 | 16.352 | 41.906 | 107.789 |
| avg mut/step | 1.730 | 1.936 | 2.918 | 2.267 | 4.355 | 8.463 | 22.369 | 70.989 | 225.029 |
| avg eval count | 2,575 | 2,732 | 2,734 | 2,776 | 2,809 | 3,165 | 3,486 | 4,389 | 6,151 |
| max eval count | 73,911 | 34,215 | 75,791 | 42,620 | 25,352 | 31,966 | 37,725 | 50,454 | 55,022 |
| min eval count | 33 | 17 | 11 | 0 | 35 | 9 | 5 | 23 | 19 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The results for the $pmut_\beta$ operator are even more clear than for the (1+1) EA. With decreasing values for β the amounts of flipped bits per step increases. The performance on the other hand decreases with rising values for β which fits into the theory of one bit flips being better for geometric distributed inputs. The only special case that does not support this theory is the value $\beta = 2.5$ performing better than $\beta = 2.75$. Here the same might hold as for the (1+1) EA. The number of repetitions of the algorithm might simply be too small to make the small difference in the performance between the two values visible. The difference in the performance for the $pmut_\beta$ operator is not as drastic as for the (1+1) EA. Only $\beta = 1.25$ performs significantly worse the next best value.

5.4.4. Comparison of the best variants

| | | | |
|----------------|-------------------------------|--------|--------|
| algo type | RLS _k ^S | pmut | EA-SM |
| algo param | s=2 | 3.25 | 2/n |
| avg mut/change | 1.483 | 1.693 | 2.258 |
| avg mut/step | 1.500 | 1.729 | 2.001 |
| avg eval count | 2,407 | 2,695 | 3,421 |
| max eval count | 23,155 | 57,661 | 52,762 |
| min eval count | 19 | 11 | 19 |
| fail ratio | 0.000 | 0.000 | 0.000 |

For the geometric distribution once again a variant of the RLS performs the best. $pmut_{3.25}$ performs almost equally good with only the (1+1) EA variant performing clearly worse. Here the algorithms are sorted again by their average number of flips per steps. The theory of flipping one bit per step being better seems to be true for this kind of input.

To further confirm the best choice for this kind of input there was another experiment of the best variants. The setup is mostly the same except for having a fixed time limit of 100,000 instead of using $100 \cdot n \ln(n)$ as the limit. The smaller inputs are harder relative to their input size so using $100 \cdot n \ln(n)$ as a bound is too small. The first try was executed with 50,000 as the step limit but there the algorithms performed too bad for $n = 20$. Therefore for the second attempt the step limit was increased to 100,000. The first table lists the number of runs where the different algorithms did not find the optimal solution within the time limit.

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|-----|-----|-----|-----|------|------|-------|
| RLS | 983 | 951 | 885 | 619 | 425 | 42 | 0 |
| RLS ₂ ^S | 881 | 578 | 171 | 0 | 0 | 0 | 0 |
| (1+1) EA (1/n) | 464 | 132 | 48 | 1 | 1 | 0 | 0 |
| (1+1) EA (2/n) | 149 | 11 | 1 | 0 | 0 | 0 | 0 |
| pmut(3.25) | 284 | 67 | 27 | 0 | 0 | 0 | 0 |
| pmut(3.5) | 312 | 91 | 38 | 0 | 1 | 0 | 0 |

For small inputs the geometric distributed input seems to have inputs without a perfect partition because there were many iterations where neither of the algorithms found an optimal solution within the time limit. It is still likely to have a perfect partition even for the small values in comparison to other distributions which follow afterwards. It seems many of the algorithms especially the variants of the RLS seem to be likely to get stuck in a local optimum. The (1+1) EA finds an optimum in most of the runs, so the geometric distributed inputs also seem to be likely to have a perfect partition for small values. They definitely are harder to solve for smaller input sizes as the binomial inputs, but they still have a perfect partition most times. The next table visualises the average number of iterations the algorithms needed for finding an optimal solution if the algorithm managed to do so.

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|-------|-------|-------|------|------|------|-------|
| RLS | 35 | 78 | 140 | 566 | 904 | 2119 | 2188 |
| RLS ₂ ^S | 357 | 2024 | 5369 | 4687 | 3945 | 2752 | 2583 |
| (1+1) EA (1/n) | 21827 | 20775 | 14583 | 9459 | 7702 | 4358 | 3924 |
| (1+1) EA (2/n) | 18211 | 11613 | 7529 | 5266 | 4359 | 3858 | 3293 |
| pmut(3.25) | 22530 | 16421 | 11312 | 5909 | 5375 | 2843 | 2246 |
| pmut(3.5) | 24202 | 17414 | 11731 | 6503 | 5216 | 2773 | 2388 |

The variants of the (1+1) EA and of the *pmut* algorithm seem to take about 20,000 iterations for $n = 20$ if they manage to find the optimal solution. They also perform better and better the bigger the input gets. This is probability caused by the many additional small values that can be used for smaller adjustments to the fitness. Also a really high value does not have as much of an effect, because there are possibly other larger values which cancel each other out, if they are in different bins. The standard (1+1) EA does not only find a perfect partition less often, it also needs more iterations on average if it does. So the (1+1) EA with $p_m = 2/n$ performs indeed better for every input size. The last table again lists the total average number of steps.

| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 |
|-------------------------------|-------|-------|-------|-------|-------|------|-------|
| RLS | 98300 | 95103 | 88516 | 62115 | 43019 | 6230 | 2188 |
| RLS ₂ ^S | 88142 | 58654 | 21551 | 4687 | 3945 | 2752 | 2583 |
| (1+1) EA (1/n) | 58099 | 31232 | 18683 | 9550 | 7795 | 4358 | 3924 |
| (1+1) EA (2/n) | 30397 | 12585 | 7622 | 5266 | 4359 | 3858 | 3293 |
| pmut(3.25) | 44531 | 22021 | 13707 | 5909 | 5375 | 2843 | 2246 |
| pmut(3.5) | 47851 | 24929 | 15086 | 6503 | 5311 | 2773 | 2388 |

The RLS is only an option if the input is large enough ($n \geq 10,000$). For smaller input sizes especially for $n \leq 100$ choosing the (1+1) EA with mutation rate $2/n$ seems like the best choice. For larger values this (1+1) EA does not find an optimal solution the fastest but is still fast enough to be a viable option. Another rather save option is *pmut*_{3.25}. This algorithm performs worse for $n \leq 100$ but is still good in comparison to the other algorithms. For $n \geq 1000$ *pmut*_{3.25} starts to outperform the best version of the (1+1) EA and almost all other researched algorithms.


5.5. Uniform distributed inputs

For the uniform distribution the default values were 1 for the lower bound and 50000 for the upper bound (exclusive). The range was limited to 50000 to reduce the time the algorithms needs to find an optimal solution. The higher the values are with too few values the more likely the input is to not have a perfect partition[BCP01]. This will cause the algorithms to always reach the limit for the number of iterations which drastically increases the time needed for the experiment. The length of the input was 50000.

5.5.1. RLS Comparison

| algo type | RLS _k ^B | RLS _k ^S | RLS _k ^S | RLS _k ^B | RLS _k ^S | RLS _k ^B | RLS |
|-----------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-----------|
| algo param | b=2 | s=3 | s=4 | b=3 | s=2 | b=4 | - |
| avg mut/change | 2.000 | 1.996 | 2.476 | 3.000 | 1.502 | 4.000 | 1.000 |
| avg mut/step | 2.000 | 2.000 | 2.500 | 3.000 | 1.500 | 4.000 | 1.000 |
| total avg count | 83,118 | 104,748 | 105,513 | 112,223 | 114,486 | 121,927 | 2,443,567 |
| avg eval count | 83,118 | 104,748 | 105,513 | 112,223 | 114,486 | 121,927 | 45,834 |
| max eval count | 778,110 | 1,453,252 | 898,974 | 1,377,471 | 915,268 | 816,633 | 485,275 |
| min eval count | 197 | 126 | 45 | 212 | 271 | 155 | 128 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.447 |
| avg fail dif | - | - | - | - | - | - | 1 |

Figure 5.20.: Distribution of a random uniform input (10000 values between 1 and 100)



The picture for the RLS variants on this type of input is not clear. There is no obvious tendency for neither of the variants. The only obvious thing is the RLS being the worst of the RLS variants again. Every variant reaches the optimal solution in every case except for the RLS which only manages for 44.7 % of the inputs. The RLS_2^B seems to be the best variant for these kinds of inputs. The next best variants are the RLS_k^S with $k = 3$ and $k = 4$ which only differ by 1 %.

5.5.2. (1+1) EA Comparison

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA |
|----------------|---------|---------|-----------|-----------|-----------|-----------|
| algo param | $3/n$ | $2/n$ | $4/n$ | $5/n$ | $10/n$ | - |
| avg mut/change | 3.102 | 2.287 | 4.014 | 4.937 | 9.924 | 1.577 |
| avg mut/step | 3.000 | 2.000 | 4.000 | 5.000 | 10.000 | 1.000 |
| avg eval count | 122,098 | 122,690 | 124,634 | 132,509 | 183,213 | 213,186 |
| max eval count | 956,375 | 920,658 | 1,128,158 | 1,457,069 | 1,298,089 | 2,509,163 |
| min eval count | 174 | 188 | 265 | 384 | 6 | 111 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The (1+1) EA seems to perform better with a lower mutation rate. The values $p_m = 2/n$ and $p_m = 3/n$ reach an optimal solution equally fast. From then on the speed of convergence decreases with increasing mutation rate. The only exception from this case is the standard (1+1) EA which performs the worst despite having the lowest mutation rate. For the uniform distributed input all variants of the (1+1) EA reach an optimal solution within the step limit as for the previous input types.

5.5.3. pmut Comparison

| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
|----------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| algo param | 2.50 | 2.00 | 2.25 | 2.75 | 1.75 | 3.00 | 1.50 | 3.25 | 1.25 |
| avg mut/change | 2.866 | 8.980 | 4.204 | 2.205 | 27.674 | 1.933 | 102.803 | 1.720 | 312.822 |
| avg mut/step | 2.932 | 10.108 | 4.559 | 2.274 | 34.643 | 1.934 | 158.163 | 1.729 | 719.965 |
| avg eval count | 117,346 | 121,090 | 121,818 | 126,467 | 128,188 | 140,882 | 142,970 | 150,311 | 193,296 |
| max eval count | 1,655,807 | 1,421,071 | 1,427,930 | 2,490,695 | 2,127,979 | 1,670,194 | 1,565,473 | 1,382,253 | 1,523,513 |
| min eval count | 61 | 186 | 76 | 130 | 357 | 155 | 113 | 226 | 13 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The optimal value for β seems to be somewhere around 2.0 to 2.5. The values next to this interval start to decrease in both directions, but 1.75 and 2.75 are still relatively close to the performance of the optimal value. The values equally wide apart from 2.25 perform equally good.

5.5.4. Comparison of the best variants

| | | | |
|----------------|-------------------------------|-----------|-----------|
| algo type | RLS _k ^B | EA-SM | pmut |
| algo param | b=2 | 3/n | 2.25 |
| avg mut/change | 2.000 | 3.109 | 4.273 |
| avg mut/step | 2.000 | 3.000 | 4.555 |
| avg eval count | 84,884 | 116,576 | 124,046 |
| max eval count | 741,833 | 1,176,762 | 1,159,541 |
| min eval count | 52 | 178 | 178 |
| fail ratio | 0.000 | 0.000 | 0.000 |

For the uniform distributed input the best variant of the RLS once again seems to perform the best. But by looking at the smaller values again this does not hold in general.

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-----|-----|-----|-----|------|------|-------|-------|
| RLS ₂ ^S | 996 | 980 | 892 | 360 | 289 | 14 | 0 | 0 |
| RLS ₃ ^S | 974 | 768 | 645 | 461 | 426 | 54 | 3 | 0 |
| RLS ₄ ^S | 902 | 621 | 541 | 460 | 431 | 53 | 4 | 0 |
| (1+1) EA (2/n) | 885 | 690 | 641 | 510 | 483 | 85 | 5 | 0 |
| (1+1) EA (3/n) | 772 | 596 | 549 | 425 | 449 | 53 | 2 | 0 |
| (1+1) EA (4/n) | 697 | 509 | 453 | 439 | 396 | 44 | 3 | 0 |
| pmut (2.5) | 840 | 738 | 696 | 553 | 520 | 86 | 10 | 0 |

The RLS variants are the most likely to get stuck in a local optimum for $n \leq 100$. The (1+1) EA variants also often do not find an optimal solution, but this happens less frequently. The more values the input has the more likely it is for any of the algorithms to find a perfect partition. Between $n = 100$ and $n = 500$ the performance of the RLS₂^B drastically increases and for $n \geq 500$ this variant of the RLS stays the best variant for the remaining input sizes. Uniform distributed inputs seem to be much less likely to have a perfect partition for the small input sizes which can be explained by the of Borgs *et. al.* coefficient [BCP01].

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|-------|-------|-------|--------|--------|--------|
| RLS ₂ ^S | 364 | 1540 | 6759 | 36358 | 37324 | 77210 | 83500 | 82738 |
| RLS ₃ ^S | 5204 | 33492 | 39795 | 40336 | 38838 | 104864 | 118393 | 106196 |
| RLS ₄ ^S | 17681 | 39895 | 38795 | 39959 | 38857 | 98693 | 108780 | 107857 |
| (1+1) EA (2/n) | 36004 | 41914 | 40131 | 42050 | 38937 | 108903 | 133264 | 122042 |
| (1+1) EA (3/n) | 32791 | 38954 | 38422 | 40617 | 41017 | 103831 | 112166 | 111402 |
| (1+1) EA (4/n) | 39665 | 39926 | 40709 | 39479 | 41889 | 99899 | 127578 | 110099 |
| pmut (2.5) | 39232 | 36918 | 37162 | 39866 | 40588 | 118671 | 144665 | 128531 |

The amount of steps needed to find an optimal solution seems to be nearly constant for every algorithm as the number of steps does not strictly increase with n but sometimes even decreases for $n \leq 1000$. This is caused by the number of steps the algorithm was given. For $n \leq 1000$ the time limit was 100,000 and for the bigger values it was $10n \ln(n)$. Interestingly enough the average runtime decreases from $n = 10000$ to $n = 50000$ for most algorithms.


| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|-------|-------|-------|--------|--------|--------|
| RLS ₂ ^S | 99601 | 98030 | 89930 | 59269 | 55438 | 82091 | 83500 | 82738 |
| RLS ₃ ^S | 97535 | 84570 | 78627 | 67841 | 64893 | 122198 | 120801 | 106196 |
| RLS ₄ ^S | 91932 | 77220 | 71907 | 67578 | 65209 | 116033 | 112029 | 107857 |
| (1+1) EA (2/n) | 92640 | 81993 | 78507 | 71604 | 68430 | 135844 | 137203 | 122042 |
| (1+1) EA (3/n) | 84676 | 75337 | 72228 | 65854 | 67500 | 120899 | 113784 | 111402 |
| (1+1) EA (4/n) | 81718 | 70503 | 67568 | 66047 | 64901 | 114241 | 129959 | 110099 |
| pmut (2.5) | 90277 | 83472 | 80897 | 73120 | 71482 | 145089 | 152429 | 128531 |

My general advice would be choosing the RLS_2^B for $n \geq 500$ and the (1+1) EA with $p_m = 4/n$ otherwise.

5.6. powerlaw distributed inputs


This distribution has mostly small values, but occasionally it also generates bigger values. The lower the parameter the higher the values get and also the amount of big values increases. For a parameter of $\beta = 2.75$ and a maximum value of 10,000 the distribution looks like in Figure 5.21.

Figure 5.21.: Distribution of a random powerlaw input with $\beta = 2.75$



For a value of $\beta = 1.25$ the distribution looks a bit different. There are less small values close to one and instead also big values even over 1000. Figure 5.22 is cropped to get a more clear view for the smaller values. The higher values mostly occurred 0 to 2 times. The highest value 9948 occurred only once.

Figure 5.22.: Distribution of a random powerlaw input with $\beta = 1.25$



5.6.1. RLS Comparison

The following table lists the results for the RLS for inputs that are chosen from a powerlaw distribution with $\beta = 2.75$.

| | | | | | | | |
|----------------|-----------|-----------|-----------|-----------|-----------|-----------|--------|
| algo type | RLS_k^B | RLS_k^B | RLS_k^S | RLS_k^S | RLS_k^B | RLS_k^S | RLS |
| algo param | b=4 | b=3 | s=4 | s=3 | b=2 | s=2 | - |
| avg mut/change | 4.000 | 3.000 | 2.470 | 1.959 | 2.000 | 1.447 | 1.000 |
| avg mut/step | 4.000 | 3.000 | 2.501 | 2.000 | 2.000 | 1.501 | 1.000 |
| avg eval count | 196 | 225 | 256 | 275 | 294 | 324 | 399 |
| max eval count | 8,523 | 8,635 | 10,613 | 10,930 | 11,210 | 12,417 | 12,562 |
| min eval count | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

Due to the many small values and all values being relatively small in general the higher mutation rates perform better. So the algorithms are ranked by the amount of bits they flip in an average step. For this parameter the input is relatively easy as all variants manage to find an optimal solution in 400 steps on average for an input size of 20,000. The next table lists the results for $\beta = 1.25$ and an input size of 10,000.

| | | | | | | | |
|-----------------|-----------|-----------|-----------|-----------|-----------|-----------|-------|
| algo type | RLS_k^S | RLS_k^S | RLS_k^B | RLS_k^B | RLS_k^S | RLS_k^B | RLS |
| algo param | r=4 | r=3 | n=3 | n=2 | r=2 | n=4 | - |
| avg mut/change | 2.445 | 1.967 | 3.000 | 2.000 | 1.486 | 4.000 | 1.000 |
| avg mut/step | 2.500 | 2.001 | 3.000 | 2.000 | 1.500 | 4.000 | 1.000 |
| total avg count | 255 | 272 | 284 | 286 | 330 | 370 | 434 |
| max eval count | 847 | 885 | 1,553 | 1,149 | 1,130 | 1,879 | 1,911 |
| min eval count | 26 | 25 | 36 | 15 | 12 | 29 | 26 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The input is still easy to solve for the RLS variants, but the concept of the more bits flipped the better does not hold any more. So the optimal algorithm parameter for an input from a specific distribution is not fixed for the complete distribution. Instead the optimal value can change even within the distribution if the parameters of the distribution change. Generally the RLS_k^S variants seem to perform good for the different parameters of the distribution especially with increasing value of k .

5.6.2. (1+1) EA Comparison

The first table again shows the results for parameter $\beta = 2.75$

| | | | | | | | | |
|----------------|--------|---------|--------|-------|-------|-------|-------|-------|
| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA |
| algo param | $50/n$ | $100/n$ | $10/n$ | $5/n$ | $4/n$ | $3/n$ | $2/n$ | - |
| avg mut/change | 49.922 | 99.873 | 10.009 | 5.053 | 4.105 | 3.156 | 2.280 | 1.534 |
| avg mut/step | 49.989 | 100.017 | 9.999 | 4.999 | 4.005 | 3.003 | 2.000 | 0.999 |
| avg eval count | 84 | 103 | 111 | 157 | 184 | 208 | 273 | 461 |
| max eval count | 1,281 | 1,488 | 1,946 | 3,030 | 3,043 | 3,283 | 4,744 | 7,036 |
| min eval count | 0 | 1 | 3 | 1 | 0 | 0 | 2 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

Here the same rule holds for the RLS to some extent. Until $p_m \leq 50/n$ the speed of convergence increases but at $p_m = 100/n$ the speed decreases again. The optimal value seems to be somewhere around $p_m = 50/n$. The (1+1) variants are generally faster than all RLS variants when comparing the maximum number of iterations. For mutation rates $3/n \leq p_m \leq 100/n$ the (1+1) EA is also faster on average. The next table shows the results for a powerlaw distribution with $\beta = 1.25$.

| | | | | | | | | |
|-----------------|-------|-------|-------|-------|-------|-------|---------|---------|
| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
| algo param | 4/n | 3/n | 2/n | 5/n | - | 10/n | 50/n | 100/n |
| avg mut/change | 3.913 | 3.061 | 2.255 | 4.794 | 1.559 | 9.489 | 49.461 | 99.590 |
| avg mut/step | 4.005 | 2.999 | 2.001 | 4.999 | 1.000 | 9.997 | 50.001 | 100.000 |
| total avg count | 243 | 245 | 291 | 299 | 509 | 1,248 | 53,071 | 99,580 |
| avg eval count | 243 | 245 | 291 | 299 | 509 | 1,248 | 53,071 | 97,933 |
| max eval count | 1,107 | 752 | 949 | 1,258 | 1,701 | 9,332 | 530,320 | 753,937 |
| min eval count | 23 | 4 | 15 | 17 | 50 | 26 | 213 | 121 |
| fails | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.002 |
| avg fail dif | - | - | - | - | - | - | - | 1 |

With this setting the optimal value is shifted to somewhere around $p_m = 4/n$. The higher mutation rates perform drastically slower with $p_m = 100/n$ being 500 times slower than the optimal value. The speed of convergence is sometimes even to slow find an optimal solution in time $10 \cdot n \ln(n)$.

5.6.3. pmut Comparison

The first table again shows the results for parameter $\beta = 2.75$

| | | | | | | | | | |
|-----------------|---------|---------|--------|--------|--------|--------|--------|--------|--------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 1.25 | 1.50 | 1.75 | 2.00 | 2.25 | 2.50 | 2.75 | 3.00 | 3.25 |
| avg mut/change | 173.565 | 62.036 | 20.819 | 7.619 | 4.175 | 2.729 | 2.191 | 1.870 | 1.672 |
| avg mut/step | 370.551 | 101.051 | 28.236 | 9.183 | 4.338 | 2.894 | 2.257 | 1.939 | 1.734 |
| total avg count | 2,090 | 2,098 | 2,137 | 2,181 | 2,209 | 2,240 | 2,252 | 2,273 | 2,289 |
| avg eval count | 110 | 118 | 156 | 200 | 229 | 260 | 272 | 293 | 309 |
| max eval count | 51,062 | 36,295 | 31,245 | 28,794 | 27,556 | 27,371 | 26,718 | 27,302 | 27,065 |
| min eval count | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| fail ratio | 0.001 | 0.001 | 0.001 | 0.001 | 0.001 | 0.001 | 0.001 | 0.001 | 0.001 |
| avg fail dif | 176 | 176 | 176 | 176 | 176 | 176 | 176 | 176 | 176 |

For pmut the same holds as for the RLS. The more bits the algorithms flips on average the better the performance on average. Surprisingly the performance in the worst runs behaves inverted. The fewer bits the algorithm flips on average the more stable the search becomes. This might be caused by the really large amount of bits flipped for the lower values.

| | | | | | | | | | |
|-----------------|--------|--------|-------|-------|---------|-------|-------|-------|-------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 1.50 | 1.75 | 2.00 | 2.25 | 1.25 | 2.50 | 2.75 | 3.00 | 3.25 |
| avg mut/change | 31.034 | 12.655 | 6.445 | 3.686 | 70.708 | 2.643 | 2.147 | 1.863 | 1.676 |
| avg mut/step | 70.574 | 22.234 | 8.730 | 4.286 | 222.822 | 2.894 | 2.287 | 1.943 | 1.731 |
| total avg count | 171 | 179 | 193 | 215 | 222 | 246 | 269 | 292 | 306 |
| max eval count | 590 | 475 | 705 | 775 | 1,207 | 705 | 894 | 1,048 | 1,003 |
| min eval count | 16 | 28 | 19 | 16 | 31 | 7 | 21 | 24 | 36 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The optimal value here seems to be somewhere around $\beta = 1.5$, so only lightly smaller in comparison to the (1+1) EA where the optimal value almost change from one side of the spectrum to the other. Here the inverted stability of the search does not occur. The variants that take longer on average tend to also take longer in their worst runs.

5.6.4. Comparison of the best variants

The first table again shows the results for parameter $\beta = 2.75$

| | | | |
|----------------|---------|--------|-----------|
| algo type | pmut | EA-SM | RLS_k^B |
| algo param | 1.25 | $50/n$ | $b=4$ |
| avg mut/change | 216.029 | 49.921 | 4.000 |
| avg mut/step | 370.312 | 50.022 | 4.000 |
| avg eval count | 56 | 84 | 181 |
| max eval count | 2,616 | 1,244 | 3,048 |
| min eval count | 0 | 1 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 |

The ranking follows the amount of bits the algorithms flip on average per step. $pmut_{1.25}$ manages to find the solution in just 56 iterations on average. The (1+1) EA with $p_m = 50/n$ is slower than $pmut_{1.25}$ but instead has a lower value for the maximum number of iterations. Both options seem fine. Even the RLS_4^B is still very fast for the powerlaw distributed input with $\beta = 2.75$. For $\beta = 1.25$ the results are a bit different.

| | | | |
|----------------|--------|-----------|-------|
| algo type | pmut | RLS_k^S | EA-SM |
| algo param | 1.50 | $r=4$ | $4/n$ |
| avg mut/change | 30.913 | 2.438 | 3.917 |
| avg mut/step | 70.577 | 2.501 | 4.000 |
| avg eval count | 178 | 251 | 251 |
| max eval count | 629 | 778 | 1,018 |
| min eval count | 17 | 19 | 13 |
| fail ratio | 0.000 | 0.000 | 0.000 |

The RLS_4^S now performs equally good as the (1+1) EA variant with $p_m = 4/n$, but is still slower than $pmut_{1.5}$. As the first inputs were less difficult to solve than the inputs with $\beta = 1.25$ the second value was chosen for the evaluation of smaller input sizes.

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|--------------------|-----|-----|-----|-----|------|------|-------|-------|
| RLS_3^S | 822 | 400 | 87 | 0 | 0 | 0 | 0 | 0 |
| RLS_4^S | 806 | 389 | 80 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA ($3/n$) | 803 | 364 | 51 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA ($4/n$) | 803 | 359 | 49 | 0 | 0 | 0 | 0 | 0 |
| pmut (1.5) | 802 | 353 | 41 | 0 | 0 | 0 | 0 | 0 |
| pmut (1.75) | 804 | 354 | 41 | 0 | 0 | 0 | 0 | 0 |

The RLS is once again the algorithm that is the most likely to be stuck in a local optimum. Compared to the other algorithms it is not as drastic as for the binomial input for example. Only for $n < 500$ the algorithms do not find a global optimum in every run. The setting of the parameter almost doesn't affect the amount of runs without an optimal result. The main differences are between the different algorithms themselves. This type of input is probably easy to solve if it has a perfect partition. The two stopping conditions where a step limit and finding a perfect partition or a partition with difference of one between the two bin for uneven n . So in 80% of the runs the algorithms might have found an optimal solution, but the stopping conditions did not trigger as the solutions where not close to a perfect partition.

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|--------------------|------|------|------|-----|------|------|-------|-------|
| RLS_3^S | 523 | 2517 | 1237 | 139 | 158 | 226 | 279 | 509 |
| RLS_4^S | 1291 | 1976 | 1935 | 141 | 151 | 203 | 246 | 426 |
| (1+1) EA ($3/n$) | 554 | 1849 | 1218 | 152 | 168 | 210 | 248 | 401 |
| (1+1) EA ($4/n$) | 382 | 1107 | 633 | 175 | 190 | 223 | 252 | 359 |
| pmut (1.5) | 667 | 487 | 164 | 147 | 148 | 163 | 178 | 206 |
| pmut (1.75) | 465 | 665 | 229 | 135 | 135 | 160 | 179 | 244 |

Looking at the time the algorithms needed on average the runs that hit the step limit could have possibly been no failed runs. The easiest are inputs with size $n = 500$. For smaller values of n the algorithms sometimes fail and even in a good run they need more iterations to find an optimal solution. Due to the increasing size of the input the algorithms need more time for the bigger values.

| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|------|-----|------|------|-------|-------|
| RLS ₃ ^S | 82293 | 41510 | 9829 | 139 | 158 | 226 | 279 | 509 |
| RLS ₄ ^S | 80850 | 40107 | 9780 | 141 | 151 | 203 | 246 | 426 |
| (1+1) EA (3/ n) | 80409 | 37576 | 6256 | 152 | 168 | 210 | 248 | 401 |
| (1+1) EA (4/ n) | 80375 | 36609 | 5502 | 175 | 190 | 223 | 252 | 359 |
| pmut (1.5) | 80332 | 35615 | 4258 | 147 | 148 | 163 | 178 | 206 |
| pmut (1.75) | 80491 | 35829 | 4320 | 135 | 135 | 160 | 179 | 244 |

$pmut_{1.75}$ is not only the best variant for the bigger values of n but also for smaller inputs as well. It is the least likely to be stuck in a local optimum, and it is also the fastest if it reaches a global optimum.

5.7. OneMax Equivalent for PARTITION

The input of this subsection is more or less equivalent to the OneMax problem. All values except the last are either 1 or follow any distribution. The last value is the sum of all other values. The optimal solution is therefore the 000...01 or the 111...01 string. So the best solution is almost identical to OneMax/ZeroMax depending on the value of the last bit. For OneMax the mutation rate of $1/n$ is proven to be optimal for the (1+1) EA [Wit13]. This should also hold for this input. The RLS variants should also perform worse than the standard RLS. The higher the value for β the better the $pmut_\beta$ mutation should perform. Flips of the first bits could decrease the runtime, depending on how often they happen. By doing some testing with various algorithm variants of the RLS and the (1+1) EA it looked like the last bit was only flipped at most once for every input. There was only one case where it was flipped twice, but it was never flipped more than twice per run. The average number of flips was also mostly closer to zero than to one.

The experiments were conducted with the variant where the smaller values are all one. So for every run of each algorithm the input was $[1, 1, \dots, 1, n - 1]$. An input like this takes less time for every algorithm, but the results are mostly the same. For some experiments not all 1000 repetitions were executed as there was a clear tendency which of the algorithms performs better.

5.7.1. RLS Comparison

| algo type | RLS | RLS _k ^S | RLS _k ^S | RLS _k ^S | RLS _k ^B | RLS _k ^B | RLS _k ^B |
|-----------------|---------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|
| algo param | - | s=2 | s=3 | s=4 | b=3 | b=2 | b=4 |
| avg mut/change | 1.000 | 1.181 | 1.688 | 1.865 | 3.000 | 1.997 | 3,997 |
| avg mut/step | 1.000 | 1.500 | 2.000 | 2.500 | 3.000 | 2.000 | 3.000 |
| total avg count | 90,931 | 168,311 | 236,317 | 307,533 | 921,030 | 921,030 | 921,030 |
| avg eval count | 90,931 | 168,311 | 236,317 | 307,533 | - | - | - |
| max eval count | 156,854 | 296,206 | 498,474 | 595,831 | - | - | - |
| min eval count | 64,941 | 120,582 | 158,304 | 212,193 | - | - | - |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 1.000 | 1.000 | 1.000 |
| avg fail dif | - | - | - | - | 36 | 53 | 263 |

As expected the standard RLS reaches an optimal solution the fastest. It also reaches the optimal value for every instance. The RLS_k^S variants need more iterations to find an optimal solution. By looking at the average values more closely it seems like the

average number of steps for the RLS_k^S is roughly $25,000 + 70,000k \pm 5,000$. The standard RLS is equivalent to RLS_k^S or RLS_k^B with $k = 1$. So the value of $k = 1$ seems to be optimal for the RLS variants too. The RLS_k^B variants on the other hand do not reach any of the two optimal solutions in any run. This is most likely caused by their very low possibility of flipping only one bit in a single step. They would eventually reach the optimal solution as well, but this would take much longer than for the RLS. The probability of flipping only one bit in a step is $\mathcal{O}(n^{1-k})$ which results in a single bit flip every $\mathcal{O}(n^{k-1})$ steps in expectation. Because the fitness can only improve for OneMax making steps flipping more bits does not harm the fitness. The bound for OneMax is $\mathcal{O}(n \log n)$ and with the previous result the expected number of steps is bounded by $\mathcal{O}(n \cdot \mathcal{O}(n^{k-1}) \cdot \log(n \cdot \mathcal{O}(n^{k-1}))) = \mathcal{O}(n^{k-1+1} \cdot (k-1+1) \cdot \log(n)) = \mathcal{O}(kn^k \cdot \log(n))$. This problem is not equivalent to OneMax, as a flip of the bit with the highest value inverts the fitness function to ZeroMax but the result might still hold as the bound for the standard RLS for this input is the same as for the RLS on OneMax.

5.7.2. (1+1) EA Comparison

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
|-----------------|---------|---------|---------|---------|---------|---------|---------|---------|
| algo param | - | $2/n$ | $3/n$ | $4/n$ | $5/n$ | $10/n$ | $50/n$ | $100/n$ |
| avg mut/change | 1.273 | 1.750 | 2.334 | 2.965 | 3.636 | 7.288 | 45.923 | 95.590 |
| avg mut/step | 1.000 | 2.000 | 3.000 | 4.000 | 5.000 | 10.000 | 49.998 | 100.00 |
| total avg count | 230,328 | 297,602 | 495,951 | 860,736 | 921,030 | 921,030 | 921,030 | 921,030 |
| avg eval count | 230,328 | 297,602 | 495,951 | 812,983 | - | - | - | - |
| max eval count | 399,393 | 625,976 | 839,325 | 917,029 | - | - | - | - |
| min eval count | 162,400 | 193,796 | 347,185 | 635,812 | - | - | - | - |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.442 | 1.000 | 1.000 | 1.000 | 1.000 |
| avg fail dif | - | - | - | 1 | 18 | 570 | 2,488 | 3,115 |

This experiment was terminated after 224 runs of the algorithms, as the results were already clear enough. For this input the same as for OneMax holds. The static mutation rate $p_m = 1/n$ is the optimal value and the performance of the (1+1) EA decreases with rising mutation rate. Only for $p_m \leq 3/n$ the (1+1) EA managed to find one of the two optimal solutions in $10 \cdot n \ln(n)$ steps every time. With mutation rate $p_m = 4/n$ the (1+1) EA only managed to find the optimal solution in about 55 % of the inputs. The remaining mutation rates did not manage to find an optimal solution in any of the runs. Another interesting fact is the average number of bits flipped in a successful step. For the other inputs the overall average number of bits flipped in any step was mostly the same as for the average value of the successful steps. Here this is not the case. All mutation rates flipped fewer bits in the successful steps than in the average step. The only exception is the standard mutation rate which is caused by the steps where the algorithm would flip no bit. Those steps decrease the number of the average case but not of the successful case as those steps were skipped.

5.7.3. pmut Comparison

| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
|-----------------|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| algo param | 3.25 | 3.00 | 2.75 | 2.50 | 2.25 | 2.00 | 1.75 | 1.50 | 1.25 |
| avg mut/change | 1.289 | 1.359 | 1.459 | 1.591 | 1.813 | 2.207 | 2.760 | 3.604 | 5.382 |
| avg mut/step | 1.731 | 1.934 | 2.270 | 2.907 | 4.371 | 8.486 | 22.299 | 70.692 | 224.466 |
| total avg count | 145,095 | 149,664 | 170,912 | 181,099 | 214,361 | 249,102 | 301,566 | 415,413 | 715,219 |
| avg eval count | 145,095 | 149,664 | 170,912 | 181,099 | 214,361 | 249,102 | 301,566 | 415,413 | 683,204 |
| max eval count | 217,932 | 223,561 | 254,330 | 246,635 | 365,378 | 376,768 | 431,629 | 735,214 | 853,181 |
| min eval count | 111,061 | 119,931 | 120,965 | 130,174 | 161,244 | 180,570 | 232,166 | 311,979 | 492,686 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.135 |
| avg fail dif | - | - | - | - | - | - | - | - | 1 |

The results for the *pmut* operator are pretty similar to the results for the (1+1) EA and the RLS. The parameter $\beta = 3.25$ which flips the least bits on average finds the solution the fastest. The other values for β increase the time needed for finding one of the two optimums with decreasing value for β . All variants find an optimum in every run except for $\beta = 1.25$ which has a much higher value for the number of flipped bits per steps. The average number for the number of bits flipped in a successful mutation is much lower than for the other inputs especially for the lower values for β . For the binomial and geometric input the successful average was around 100 for $\beta = 1.25$ but for the OneMax equivalent it was only at 5.

5.7.4. Comparison of the best variants

| | | | |
|----------------|---------|---------|---------|
| algo type | RLS | pmut | EA |
| algo param | - | 3.25 | - |
| avg mut/change | 1.000 | 1.287 | 1.272 |
| avg mut/step | 1.000 | 1.729 | 1.000 |
| avg eval count | 91,171 | 143,121 | 231,082 |
| max eval count | 153,143 | 227,737 | 446,942 |
| min eval count | 65,783 | 93,602 | 165,818 |
| fail ratio | 0.000 | 0.000 | 0.000 |

The results for this experiment are as expected. All three algorithms find the optimal value within the time limit. The RLS performs better than the (1+1) EA because it does only single bit flips. The *pmut*_{3.25} performs better than the standard (1+1) EA although flipping more bits on average. This is most likely cause by the few steps where *pmut* flips many bits which increase the average. But *pmut* most likely chooses to flip only one bit more often as the (1+1) EA.



For this comparison neither of the algorithms failed to find one of the two optimal solutions. The following table lists the amount of iterations the algorithms needed to find an optimal solution.

| | | | | | | | | |
|----------------|-----|-----|------|------|-------|--------|--------|---------|
| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS | 56 | 187 | 446 | 3050 | 6725 | 42056 | 90485 | 537292 |
| (1+1) EA (1/n) | 131 | 438 | 1059 | 7435 | 16963 | 106410 | 230310 | 1379632 |
| pmut (3.25) | 82 | 277 | 667 | 4666 | 10432 | 66450 | 142919 | 846722 |
| pmut (3.5) | 77 | 265 | 642 | 4393 | 9930 | 62337 | 137048 | 801570 |

The RLS performs the best closely follow by both *pmut* variants. The standard (1+1) EA performs a bit worse than the other three algorithms and approaches $en \ln(n)$ instead of staying close to $n \ln(n)$.

In a previous chapter the $\mathcal{O}(n \log n)$ bound was proven for the (1+1) EA and the RLS (Theorem 3.1). This seems to hold in practice at least for the easiest version of this input where the small values are one (see figure 5.23). Another variant of this input are uniform distributed inputs for example. The small values in this case are chosen from $\sim U(1, 49999)$ and the last value again is the sum of all other values. This input is harder because switching multiple small values for a big value increases the fitness but increases the Hamming distance to the optimum.

Looking at figure 5.24 the uniform distributed variant of the OneMax equivalent looks not much harder compared to the variant where all values are one. The graphs are almost identical. It was clear for the RLS because the RLS can't switch elements and there behaves exactly the same. But even the other algorithms that are able to switch seem to not do it very often as the runtime is almost the same.

Figure 5.23.: Runtime for OneMax equivalent with a $n \ln(n)$ scale

Figure 5.24.: Runtime for OneMax equivalent with uniform distribution on a $n \ln(n)$ scale


5.8. Carsten Witt's worst case input

C. Witt proved that the RLS and the (1+1) EA find a $(4/3 + \epsilon)$ approximation in expected time $\mathcal{O}(n)$ and a $(4/3)$ -approximation in expected time $\mathcal{O}(n^2)$ [Die05].

TODO: why is this citation wrong He then introduced an almost worst case input to prove the bound for the approximation ratio is at least almost tight. The input is defined as followed for any $0 < \epsilon < 1/3$ and even n :

The input contains two numbers of value $1/3 - \epsilon/4$ and $n - 2$ elements of value $(1/3 + \epsilon/2)/(n - 2)$. The total volume is normalised to 1. When the two large values are in the same bin, the RSHs are tricked into a local optima, where only w_1 and w_2 are in the first bin and the remaining elements in the other bin. This results in an almost worst case. As all researched inputs in this paper contained only integer values this input is adjusted a bit. To prevent the small values to be below zero they are instead normalised to 1. The two big values are scaled by the same factor of $((1/3 + \epsilon/2)/(n - 2))^{-1}$. The higher the value for ϵ the more likely the input is to get stuck in the local optima. With increasing ϵ the local optima becomes better and better. For the small values of ϵ there were only a few cases where some algorithms did not find an optimal solution. To make this effect more visible the value of ϵ was set to $\epsilon = 0.3$.

For $n = 10,000$ this evaluates to $w_1 = w_2 = 5344$ and $W = 9998 \cdot 1 + 2 \cdot 5344 = 20686$. The input then looks like this: $[1, 1, \dots, 1, 1, 5344, 5344]$. The fitness of the local optimum is $f(x) = 2 \cdot 5433 = 10688$. To leave the local optimum the algorithm therefore has to

flip at least $5433 + 9998 - 10688 = 4654$ bits as well in the same step. The best fitness is $f(x) = 5344 + 9998/2 = 10343$, which leads to a difference of $f(localOptimum) - f(opt) = 345$ and a approximation ratio of $f(localOptimum)/f(opt) = 10688/10343 = 1.033$. This is not really close to the worst case of $4/3$ any more but with this setting at least many algorithms are stuck in the local optimum at least once for the 1000 runs.

5.8.1. RLS Comparison

| algo type | RLS _k ^B | RLS _k ^B | RLS _k ^B | RLS _k ^S | RLS _k ^S | RLS _k ^S | RLS |
|-----------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|-------------------------------|---------|
| algo param | b=4 | b=3 | b=2 | s=4 | s=3 | s=2 | - |
| avg mut/change | 3.997 | 3.000 | 1.998 | 2.344 | 1.967 | 1.307 | 1.000 |
| avg mut/step | 4.000 | 3.000 | 2.000 | 2.501 | 2.000 | 1.500 | 1.000 |
| total avg count | 4,587 | 4,844 | 9,607 | 11,503 | 12,841 | 29,741 | 114,452 |
| avg eval count | 4,587 | 1,165 | 6,864 | 1,387 | 1,810 | 2,176 | 2,376 |
| max eval count | 56,029 | 10,171 | 75,524 | 11,060 | 11,544 | 12,930 | 12,023 |
| min eval count | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| fail ratio | 0.000 | 0.004 | 0.003 | 0.011 | 0.012 | 0.030 | 0.122 |
| avg fail dif | - | 382 | 401 | 345 | 345 | 345 | 345 |

The RLS is by far most likely to get stuck in the local optimum. The general tendency is the more bits the algorithm flips in expectation the more unlikely the local optimum becomes. The only case where this does not hold is the RLS₂^B being better than the RLS₄^S, although the RLS₄^S flips more bits in expectation. This might be caused by the higher probability for flipping only one bit for the RLS₄^S. This causes the algorithm to find the local optimum faster before separating w_1 and w_2 . So the ranking of the algorithms is completely inverted compared to the OneMax input. The RLS₂^B and RLS₃^B both had runs where they neither found the global nor one of the two local optima. The algorithms were most likely tricked into the direction of the local optimum and did not manage to leave it. But they were also not fast enough to reach the local optimum because of their low probability to flip only one bit.

5.8.2. (1+1) EA Comparison

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA |
|-----------------|---------|--------|--------|-------|--------|--------|--------|--------|
| algo param | 100/n | 50/n | 10/n | 5/n | 4/n | 3/n | 2/n | - |
| avg mut/change | 100.074 | 49.998 | 10.024 | 5.083 | 4.110 | 3.130 | 2.224 | 1.438 |
| avg mut/step | 100.095 | 50.024 | 9.995 | 5.001 | 4.002 | 2.999 | 2.000 | 1.000 |
| total avg count | 64 | 106 | 414 | 855 | 1,024 | 4,980 | 9,252 | 38,213 |
| avg eval count | 64 | 106 | 414 | 855 | 1,024 | 1,301 | 1,899 | 3,341 |
| max eval count | 407 | 898 | 4,026 | 9,068 | 10,830 | 10,540 | 13,135 | 19,361 |
| min eval count | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.004 | 0.008 | 0.038 |
| avg fail dif | - | - | - | - | - | 345 | 345 | 345 |

For the (1+1) EA the result is also the inversion of the results for the OneMax equivalent. The higher the mutation rate the faster the algorithms reach a global optimum. This holds at least up to $p_m \leq 100/n$. With mutation rate $p_m \leq 3/n$ the algorithm reaches the worst case at least once in 1000 runs. If the algorithm did not manage to find an optimal solution the fitness was always the same. So there was no run where any algorithm neither found a global nor the local optimum.

5.8.3. pmut Comparison

| | | | | | | | | | |
|-----------------|---------|--------|--------|-------|-------|--------|--------|--------|--------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 1.25 | 1.50 | 1.75 | 2.00 | 2.25 | 2.50 | 2.75 | 3.00 | 3.25 |
| avg mut/change | 192.353 | 68.078 | 23.956 | 8.809 | 4.373 | 2.860 | 2.111 | 1.770 | 1.578 |
| avg mut/step | 220.467 | 68.897 | 22.423 | 8.578 | 4.410 | 2.922 | 2.271 | 1.934 | 1.729 |
| total avg count | 41 | 88 | 211 | 493 | 956 | 2,272 | 13,622 | 16,503 | 20,333 |
| avg eval count | 41 | 88 | 211 | 493 | 956 | 1,353 | 1,670 | 1,795 | 1,952 |
| max eval count | 212 | 753 | 1,704 | 5,933 | 8,957 | 19,703 | 11,437 | 11,593 | 12,032 |
| min eval count | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.001 | 0.013 | 0.016 | 0.020 |
| avg fail dif | - | - | - | - | - | 345 | 345 | 345 | 345 |

For *pmut* the result is the exact same as for the (1+1) EA. The lower β the better the performance as more bits are flipped in each step. The algorithm is tricked into the local optima only for $\beta \leq 2.5$. If the algorithm is on the path to the local optimum it is always fast enough to reach it within the time limit.

5.8.4. Comparison of the best variants

| | | | |
|----------------|---------|---------|-----------|
| algo type | pmut | EA-SM | RLS_k^B |
| algo param | 1.25 | $100/n$ | $b=4$ |
| avg mut/change | 195.454 | 99.875 | 3.998 |
| avg mut/step | 224.523 | 99.980 | 4.000 |
| avg eval count | 43 | 67 | 4,458 |
| max eval count | 240 | 521 | 43,317 |
| min eval count | 0 | 0 | 0 |
| fail ratio | 0.000 | 0.000 | 0.000 |

The $pmut_{1.25}$ and the (1+1) EA with $p_m = 100/n$ perform the best and always find an optimal solution within 600 iterations and even under 100 on the average case. The RLS_4^B performs significantly worse. For the OneMax input $pmut_{1.25}$ flipped 224 bits on average per step, but the average of the successful steps was only 5. Here the average of all steps is again 224, but the average of the successful steps is at 194. The heavy tail here really increases the performance as most of the high values are accepted. In the experiment with different input sizes the mutation rate of $p_m = 100/n$ is ≥ 1 for $n \leq 100$. If the algorithm flips every bit then it won't change its solution. In these cases the mutation rate was then set to $p_m = 1/2$.

| | | | | | | | | |
|----------------------|----|----|-----|-----|------|------|-------|-------|
| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS_4^B | 4 | 0 | 3 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA ($100/n$) | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| pmut (1.25) | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Only the RLS variant had runs where it did not reach a global optimum. This happened in less than 0.5 % of the inputs for $n = 20$ and $n = 100$. For the other input sizes it also managed to reach a global optimum for all inputs.

| | | | | | | | | |
|----------------------|----|----|-----|-----|------|------|-------|-------|
| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS_4^B | 12 | 26 | 53 | 221 | 442 | 2166 | 4194 | 20690 |
| (1+1) EA ($100/n$) | 10 | 16 | 24 | 32 | 34 | 47 | 69 | 222 |
| pmut (1.25) | 7 | 10 | 11 | 19 | 23 | 35 | 43 | 63 |

For the lower input sizes the RLS is slower than the remaining algorithm even it manages to find a global optimum.


| | | | | | | | | |
|----------------------|-----|----|-----|-----|------|------|-------|-------|
| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS_4^B | 412 | 26 | 353 | 221 | 442 | 2166 | 4194 | 20690 |
| (1+1) EA ($100/n$) | 10 | 16 | 24 | 32 | 34 | 47 | 69 | 222 |
| pmut (1.25) | 7 | 10 | 11 | 19 | 23 | 35 | 43 | 63 |

The $pmut_{1.25}$ is the best variant closely followed by the (1+1) EA. The RLS version is by far slower than the other to variants for the bigger input sizes. Even for the smaller inputs it is still slower.

5.9. Multiple distributions mixed

This input does not follow a specific distribution but rather is a mix of the previous distributions. The value is either chosen uniform random, from a binomial distribution, from a geometric distribution or a powerlaw distribution. One of the distributions is chosen uniform randomly. This process is repeated n times. The values of this distribution then follow neither of the used distributions.

Figure 5.25.: Distribution of a mixed input with $\sim U(1, 999)$, $\sim B(1000, 0.1)$, $\sim Geo(0.01)$, $\sim D_{1000}^{1.25}$



A possible distribution is shown in Figure 5.25. The spike in the curve is caused by the binomial distribution with an expected value of 100. Each value occurs at least 2.5 times in expectation due to the uniform distribution. The large spike for the lowest values is caused by both the geometric and the powerlaw distribution. The parameter for this figure were lowered to improve the visibility of bigger values which occurs much less frequently than the small values. With a larger span of possible values the big values would be even less visible.

The used distributions for the experiments in the next subsections were $\sim U(1, 49999)$, $\sim B(10000, 0.1)$, $\sim Geo(0.001)$, $\sim D_{50,000}^{1.25}$.

5.9.1. RLS Comparison

| algo type | RLS | RLS_k^S | RLS_k^S | RLS_k^S | RLS_k^B | RLS_k^B | RLS_k^B |
|----------------|-------|-----------|-----------|-----------|-----------|-----------|-----------|
| algo param | - | s=2 | s=3 | s=4 | b=2 | b=3 | b=4 |
| avg mut/change | 1.000 | 1.466 | 1.907 | 2.323 | 2.000 | 3.000 | 3.999 |
| avg mut/step | 1.000 | 1.499 | 1.999 | 2.501 | 2.000 | 3.000 | 4.000 |
| avg eval count | 340 | 361 | 414 | 479 | 835 | 2,917 | 6,847 |
| max eval count | 1,125 | 1,299 | 1,579 | 1,982 | 5,315 | 20,525 | 60,444 |
| min eval count | 21 | 30 | 32 | 47 | 66 | 42 | 62 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The results are mostly the same as for the geometric input. The performance increases with the probability of flipping only one bit. The main difference is the RLS being the

best variant in this case as it reaches an optimal solution in every run. The penalty of flipping only one bit is greater than for the geometric input but certainly not as drastic as for the OneMax equivalent. By comparing the average number of iterations it seems like this input is easier than the geometric input.

5.9.2. (1+1) EA Comparison

| | | | | | | |
|----------------|-------|-------|-------|-------|--------|---------|
| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
| algo param | $2/n$ | - | $3/n$ | $4/n$ | $5/n$ | $10/n$ |
| avg mut/change | 2.152 | 1.522 | 2.871 | 3.644 | 4.474 | 9.165 |
| avg mut/step | 1.999 | 1.001 | 2.997 | 3.998 | 4.997 | 10.000 |
| avg eval count | 494 | 545 | 653 | 1,047 | 1,787 | 27,486 |
| max eval count | 1,848 | 1,703 | 2,720 | 4,291 | 12,437 | 258,079 |
| min eval count | 33 | 64 | 31 | 28 | 27 | 81 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

For the (1+1) EA the same holds. It performs better with decreasing mutation rate with $p_m = 2/n$ being the only exception again. The same was true for the geometric input. The penalty for the wrong parameter is also bigger for the (1+1) EA compared to the geometric input. So the results are rather similar to the RLS.

5.9.3. pmut Comparison

| | | | | | | | | | |
|----------------|-------|-------|-------|-------|-------|-------|--------|--------|---------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 3.00 | 3.25 | 2.75 | 2.50 | 2.25 | 2.00 | 1.75 | 1.50 | 1.25 |
| avg mut/change | 1.814 | 1.648 | 2.114 | 2.512 | 3.351 | 6.607 | 12.646 | 32.530 | 73.371 |
| avg mut/step | 1.938 | 1.730 | 2.299 | 2.939 | 4.354 | 8.577 | 22.003 | 70.569 | 225.182 |
| avg eval count | 341 | 347 | 350 | 355 | 383 | 415 | 475 | 614 | 950 |
| max eval count | 1,338 | 1,093 | 1,291 | 1,251 | 1,398 | 1,812 | 2,925 | 3,066 | 5,350 |
| min eval count | 37 | 23 | 19 | 9 | 46 | 32 | 25 | 47 | 42 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

$pmut_\beta$ also performs similar to the geometric input. Here $\beta = 3.25$ and $\beta = 3.0$ are switched instead of $\beta = 2.5$ and 3.75 . Apart from that the algorithms perform strictly ranked by their probability to flip only one bit per step. Solving the mixed input is easier for the $pmut$ operator too. All variants manage to find an optimal solution within 1000 steps on average compared to 6,200 steps for the geometric input. The maximum number of steps is also lower by a factor of at least 10 for every value of β .

5.9.4. Comparison of the best variants

| | | | |
|----------------|-------|-------|-------|
| algo type | RLS | pmut | EA-SM |
| algo param | - | 3.00 | $2/n$ |
| avg mut/change | 1.000 | 1.804 | 2.144 |
| avg mut/step | 1.000 | 1.931 | 2.000 |
| avg eval count | 341 | 342 | 501 |
| max eval count | 942 | 1,030 | 2,280 |
| min eval count | 53 | 35 | 43 |
| fail ratio | 0.000 | 0.000 | 0.000 |

This input seems generally easy to solve as for every base algorithm multiple variants reach an optimal solution within 1000 steps. The standard RLS reaches an optimal solution the fastest, but the other algorithms are almost equally fast. All algorithms finish within 501 steps on average and always in less than 2500 steps.

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-----|-----|-----|-----|------|------|-------|-------|
| RLS | 984 | 773 | 411 | 1 | 0 | 0 | 0 | 0 |
| RLS ₂ ^S | 890 | 241 | 14 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA (1/ <i>n</i>) | 711 | 75 | 5 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA (2/ <i>n</i>) | 541 | 14 | 0 | 0 | 0 | 0 | 0 | 0 |
| pmut (3.0) | 566 | 63 | 4 | 0 | 0 | 0 | 0 | 0 |
| pmut (3.25) | 587 | 63 | 7 | 0 | 0 | 0 | 0 | 0 |

This input is only hard to solve for $n < 100$. Only the RLS is stuck in a local optimum for many inputs. The other algorithms manage to escape the local optima in most cases even for $n = 100$. For $n \geq 500$ every input is solved by each of the chosen algorithms except for the standard RLS failing once for $n = 500$. This is probably caused by the many small values from the powerlaw and geometric distribution. The longer the input the more of these small values can be used to fill the gaps.

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|------|------|-----|------|------|-------|-------|
| RLS | 71 | 145 | 198 | 222 | 232 | 288 | 337 | 534 |
| RLS ₂ ^S | 386 | 993 | 1020 | 300 | 289 | 325 | 364 | 510 |
| (1+1) EA (1/ <i>n</i>) | 15614 | 7018 | 2062 | 431 | 421 | 492 | 556 | 782 |
| (1+1) EA (2/ <i>n</i>) | 15273 | 5190 | 1295 | 457 | 432 | 478 | 507 | 636 |
| pmut (3.0) | 17101 | 5633 | 1475 | 311 | 274 | 319 | 345 | 474 |
| pmut (3.25) | 17918 | 7228 | 1691 | 283 | 268 | 313 | 332 | 475 |

The RLS variants are only able to solve the inputs if they are lucky. The (1+1) EAs and the *pmut* variants find optimal solutions more often but need way more steps if they do so. The input becomes drastically easier until $n = 500$ and very slowly becomes harder with rising n again.

| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|-------|-----|------|------|-------|-------|
| RLS | 98401 | 77332 | 41216 | 322 | 232 | 288 | 337 | 534 |
| RLS ₂ ^S | 89042 | 24854 | 2406 | 300 | 289 | 325 | 364 | 510 |
| (1+1) EA (1/ <i>n</i>) | 75612 | 13991 | 2552 | 431 | 421 | 492 | 556 | 782 |
| (1+1) EA (2/ <i>n</i>) | 61110 | 6517 | 1295 | 457 | 432 | 478 | 507 | 636 |
| pmut (3.0) | 64022 | 11578 | 1869 | 311 | 274 | 319 | 345 | 474 |
| pmut (3.25) | 66100 | 13072 | 2380 | 283 | 268 | 313 | 332 | 475 |


For $n \leq 100$ the (1+1) EA with $p_m = 2/n$ performs the best. It is still good for the bigger values of n but after $n \geq 500$ *pmut*_{3.0} reaches an optimal solution faster. The standard RLS is only the fastest for a small range of values but has the disadvantage of the poor performance for small inputs. Choosing one of the other variants is a much safer choice.

5.10. Adding multiple distributions

This input is similar to the mixed input from the last subsection. For this distribution the values are not chosen uniform random from one of the base distributions, but instead a value from every distribution is chosen and added together. This distribution will be referred to as overlapped distribution.

Figure 5.26 looks completely different from figure 5.25. No value is generated more than 20 times as opposed to the maximum amount of 350 for the mixed distribution. In this figure no distribution is clearly visible.

The used distributions for the performance experiments were $\sim U(1, 49999)$, $\sim B(10000, 0.1)$, $\sim Geo(0.001)$, $\sim D_{50000}^{1.25}$.

Figure 5.26.: Distribution of an overlapped input with $\sim U(1,999)$, $\sim B(1000,0.1)$, $\sim Geo(0.01)$, $\sim D_{1000}^{1.25}$ 

5.10.1. RLS Comparison

| algo type | RLS_k^B | RLS_k^B | RLS_k^B | RLS_k^S | RLS_k^S | RLS_k^S | RLS |
|-----------------|-----------|-----------|-----------|-----------|-----------|-----------|---------|
| algo param | b=2 | b=3 | b=4 | s=4 | s=3 | s=2 | - |
| avg mut/change | 2.000 | 3.000 | 3.999 | 2.552 | 2.064 | 1.570 | 1.000 |
| avg mut/step | 2.000 | 3.000 | 4.000 | 2.501 | 2.000 | 1.500 | 1.000 |
| total avg count | 91,899 | 123,301 | 132,017 | 153,639 | 155,128 | 179,272 | 920,109 |
| avg eval count | 91,899 | 123,301 | 131,228 | 150,557 | 152,052 | 170,263 | 199 |
| max eval count | 780,728 | 803,183 | 920,909 | 857,395 | 802,573 | 905,891 | 199 |
| min eval count | 125 | 83 | 195 | 207 | 107 | 49 | 199 |
| fail ratio | 0.000 | 0.000 | 0.001 | 0.004 | 0.004 | 0.012 | 0.999 |
| avg fail dif | - | - | 1 | 1 | 1 | 1 | 380 |

The results for this distribution are somewhat similar to the results of the binomial distribution. The best and worst algorithms are the same. Only the RLS almost fails to find an optimal solution for almost every input. For the binomial distribution the RLS_3^B also failed around 25% of the inputs but here it does not. It is even the second-best variant instead of being the second-worst. Apart from that the ranking in between is also different. The RLS_k^B variants perform the best followed by the RLS_k^S variants and lastly the RLS. For the RLS_k^B the performance decreases with increasing k but for the RLS_k^S this relation is inverted. Another big difference is the number of steps needed to find a global optimum. For the binomial input all algorithms except the RLS were really fast and only needed below 1000 iterations on average. All algorithms need around 250 times the number of steps on average for the overlapped input. This caused some of them to not be fast enough to solve the input always in time $10n \ln(n)$.

5.10.2. (1+1) EA Comparison

| | | | | | | |
|-----------------|---------|---------|---------|---------|---------|---------|
| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA |
| algo param | $4/n$ | $3/n$ | $5/n$ | $2/n$ | $10/n$ | - |
| avg mut/change | 4.023 | 3.150 | 4.914 | 2.364 | 9.790 | 1.680 |
| avg mut/step | 4.000 | 3.000 | 5.000 | 2.000 | 10.000 | 1.000 |
| total avg count | 138,106 | 151,488 | 162,312 | 187,339 | 206,040 | 328,545 |
| avg eval count | 136,537 | 146,843 | 156,964 | 177,676 | 195,888 | 266,350 |
| max eval count | 886,825 | 903,598 | 882,645 | 918,488 | 905,911 | 913,735 |
| min eval count | 237 | 330 | 196 | 218 | 202 | 608 |
| fail ratio | 0.002 | 0.006 | 0.007 | 0.013 | 0.014 | 0.095 |
| avg fail dif | 1 | 1 | 1 | 1 | 1 | 1 |

The results for the (1+1) EA are similar but not the same. Algorithms that had a good runtime on binomial inputs also have a good runtime in the overlapped runs, but the ranking is not the same. For this input the 1st and 2nd place switched, the same for 3rd and 4th and also for 5th and 6th place. This may be only caused by chance and the performance within a pair being really close but might also be caused by something else. The (1+1) EA variants also needed about 200 times as long as for the binomial inputs. All variants had at least 2 runs where they did not find an optimal solution. These runs all had a remaining difference of one to the optimal value. More time would probably be enough to reach an optimal solution in most cases.

5.10.3. pmut Comparison

| | | | | | | | |
|-----------------|---------|---------|---------|---------|---------|---------|---------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 1.75 | 2.00 | 2.25 | 2.75 | 2.50 | 1.50 | 3.00 |
| avg mut/change | 13.944 | 6.833 | 3.795 | 2.277 | 2.806 | 41.097 | 2.021 |
| avg mut/step | 22.231 | 8.485 | 4.362 | 2.270 | 2.905 | 70.686 | 1.934 |
| total avg count | 178,149 | 185,178 | 192,095 | 194,246 | 197,694 | 201,656 | 222,344 |
| avg eval count | 169,887 | 179,990 | 179,489 | 181,677 | 185,933 | 185,474 | 205,163 |
| max eval count | 910,588 | 920,835 | 908,823 | 909,803 | 907,001 | 920,331 | 904,854 |
| min eval count | 85 | 79 | 59 | 15 | 387 | 68 | 70 |
| fail ratio | 0.011 | 0.007 | 0.017 | 0.017 | 0.016 | 0.022 | 0.024 |
| avg fail dif | 1 | 1 | 1 | 1 | 1 | 1 | 1 |

For *pmut* the results are mostly the same as for the (1+1) EA. The performance is about 250 times worse than for the binomial inputs, but the ranking is still pretty close. Here the optimal parameter is around $\beta = 1.75$ instead of $\beta = 2.25$. For *pmut* the step limit was too small as well, but the difference to the optimal solution was only one on average.

5.10.4. Comparison of the best variants

| | | | |
|-----------------|-----------|---------|---------|
| algo type | RLS_k^B | EA-SM | pmut |
| algo param | b=2 | $4/n$ | 1.75 |
| avg mut/change | 2.000 | 4.002 | 15.610 |
| avg mut/step | 2.000 | 4.000 | 22.255 |
| total avg count | 90,152 | 151,282 | 180,435 |
| avg eval count | 90,152 | 146,636 | 175,214 |
| max eval count | 578,840 | 920,620 | 914,183 |
| min eval count | 242 | 54 | 153 |
| fail ratio | 0.000 | 0.006 | 0.007 |
| avg fail dif | - | 1 | 1 |

The results here are the same as for the binomial input. The RLS_2^B performs better than the (1+1) EA and $pmut_\beta$ mutation for all values of c/n and β by a factor of 1.5 with a step

limit of $10 \cdot n \ln(n)$. Due to its higher speed of convergence it is also the only algorithm to find a global optimum in every run. For the lower values of n this does not hold.

| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-----|-----|-----|-----|------|------|-------|-------|
| RLS ₂ ^S | 994 | 981 | 899 | 374 | 340 | 21 | 2 | 0 |
| RLS ₃ ^B | 976 | 767 | 467 | 426 | 423 | 51 | 4 | 0 |
| (1+1) EA (3/ n) | 807 | 605 | 578 | 493 | 503 | 62 | 2 | 0 |
| (1+1) EA (4/ n) | 705 | 550 | 499 | 484 | 468 | 68 | 4 | 0 |
| pmut (1.75) | 788 | 684 | 657 | 544 | 540 | 115 | 18 | 0 |
| pmut (2.0) | 844 | 734 | 697 | 559 | 546 | 104 | 12 | 0 |

Here almost no algorithm manages to find an optimal solution in most cases as all algorithms fail for more than 70 % of the inputs. The RLS variants perform the worst again for the small input sizes. Only for $n \geq 500$ the RLS₂^B is a good option due to the huge performance increase between $n = 100$ and $n = 500$. Overlapped inputs stay hard to solve relatively long as only for $n \geq 50,000$ all best variants of each algorithm find an optimal solution for every input.

| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|-------|-------|-------|--------|--------|--------|
| RLS ₂ ^S | 314 | 2471 | 7496 | 36412 | 38128 | 84780 | 88098 | 85872 |
| RLS ₃ ^B | 2324 | 23611 | 39185 | 41354 | 41587 | 106181 | 120997 | 124832 |
| (1+1) EA (3/ n) | 35981 | 40501 | 42587 | 42726 | 41801 | 112273 | 145284 | 151792 |
| (1+1) EA (4/ n) | 38418 | 39769 | 41415 | 43110 | 42387 | 107131 | 152592 | 154707 |
| pmut (1.75) | 42421 | 44267 | 39775 | 41363 | 42002 | 121456 | 162158 | 184272 |
| pmut (2.0) | 39690 | 41358 | 35831 | 42948 | 43160 | 126116 | 169246 | 181876 |

For $n \leq 1000$ the performance seems almost constant, but this is caused by the constant step limit of 100,000. After $n = 1000$ the value of $10 \cdot n \ln(n)$ was bigger than 100,000.


| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
|-------------------------------|-------|-------|-------|-------|-------|--------|--------|--------|
| RLS ₂ ^S | 99401 | 98146 | 90657 | 60194 | 59164 | 91942 | 89764 | 85872 |
| RLS ₃ ^B | 97655 | 82201 | 67585 | 66337 | 66295 | 122484 | 124197 | 124832 |
| (1+1) EA (3/ n) | 87644 | 76498 | 75771 | 70962 | 71075 | 131716 | 146835 | 151792 |
| (1+1) EA (4/ n) | 81833 | 72896 | 70649 | 70644 | 69350 | 128804 | 155665 | 154707 |
| pmut (1.75) | 87793 | 82388 | 79342 | 73261 | 73321 | 156462 | 175817 | 184272 |
| pmut (2.0) | 90591 | 84401 | 80556 | 74840 | 74194 | 157289 | 178268 | 181876 |

No algorithm is very successful for the small values of n . Choosing the (1+1) EA for $n < 500$ should result in the best runtime in most of the cases. Only the RLS₄^B is faster for $50 \leq n \leq 100$. For bigger input sizes the RLS₂^B is clearly the best option.

5.11. Multiple distributions mixed & overlapped

This distribution is again similar to the mixed distribution. With probability 1/2 the value is chosen from the overlapped distribution and with remaining probability 1/2 the value is chosen from the mixed distribution.

The used distributions were $\sim U(1, 49999)$, $\sim B(10000, 0.1)$, $\sim Geo(0.001)$, $\sim D_{50000}^{1.25}$. By looking at Figure 5.27 it looks like the mixed and overlapped distribution is closer to the mixed distribution than to the overlapped distribution. The results should also be closer to the mixed distribution.

Figure 5.27.: Distribution of a mixed and overlapped input with $\sim U(1, 999)$, $\sim B(1000, 0.1)$, $\sim Geo(0.01)$, $\sim D_{1000}^{1.25}$ 

5.11.1. RLS Comparison

| algo type | RLS | RLS_k^S | RLS_k^S | RLS_k^S | RLS_k^B | RLS_k^B | RLS_k^B |
|----------------|-------|-----------|-----------|-----------|-----------|-----------|-----------|
| algo param | - | s=2 | s=3 | s=4 | b=2 | b=3 | b=4 |
| avg mut/change | 1.000 | 1.435 | 1.833 | 2.221 | 2.000 | 2.999 | 4.000 |
| avg mut/step | 1.000 | 1.501 | 1.999 | 2.501 | 2.000 | 3.000 | 4.000 |
| avg eval count | 432 | 598 | 830 | 1,073 | 2,831 | 17,483 | 55,532 |
| max eval count | 1,534 | 2,915 | 6,576 | 4,947 | 18,933 | 137,135 | 504,436 |
| min eval count | 40 | 35 | 81 | 51 | 52 | 103 | 148 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

As expected the results are similar to the mixed input. There is a clear preference of algorithms with higher probability to flip only one bit per step. Mixed and geometric distributions did not lead to very bad performance for the RLS_4^B . The mixed and overlapped input punishes the worse mutation operators more than the other two inputs. In that sense this type of input is closer to the OneMax equivalent, but still way less extreme. Here still every variant reaches an optimal solution in every case.

5.11.2. (1+1) EA Comparison

| algo type | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM | EA-SM |
|----------------|-------|-------|-------|--------|--------|---------|
| algo param | - | $2/n$ | $3/n$ | $4/n$ | $5/n$ | $10/n$ |
| avg mut/change | 1.479 | 2.057 | 2.720 | 3.444 | 4.293 | 9.397 |
| avg mut/step | 1.000 | 1.999 | 3.001 | 4.001 | 5.001 | 10.000 |
| avg eval count | 892 | 991 | 1,526 | 2,708 | 4,980 | 91,623 |
| max eval count | 3,588 | 4,475 | 6,059 | 13,334 | 25,079 | 721,020 |
| min eval count | 40 | 78 | 99 | 117 | 77 | 84 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

The results here are pretty similar to the results of the RLS. The lower mutation rates perform better and higher mutation rates start to perform worse in comparison to the mixed input.

5.11.3. pmut Comparison

| | | | | | | | | | |
|----------------|-------|-------|-------|-------|-------|-------|--------|--------|---------|
| algo type | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut | pmut |
| algo param | 3.25 | 3.00 | 2.75 | 2.50 | 2.25 | 2.00 | 1.75 | 1.50 | 1.25 |
| avg mut/change | 1.604 | 1.751 | 1.977 | 2.424 | 3.266 | 5.140 | 13.276 | 31.477 | 76.065 |
| avg mut/step | 1.731 | 1.935 | 2.274 | 2.901 | 4.382 | 8.463 | 22.347 | 70.468 | 224.657 |
| avg eval count | 574 | 606 | 616 | 670 | 740 | 860 | 1,012 | 1,297 | 2,201 |
| max eval count | 1,990 | 2,304 | 2,109 | 2,625 | 2,968 | 3,283 | 4,888 | 6,196 | 10,991 |
| min eval count | 32 | 70 | 47 | 55 | 67 | 65 | 33 | 88 | 64 |
| fail ratio | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 | 0.000 |

Here the same holds. The higher mutation rates are less impacted than the higher rates for the RLS and the (1+1) EA.

5.11.4. Comparison of the best variants

| | | | |
|----------------|-------|-------|-------|
| algo type | RLS | pmut | EA |
| algo param | - | 3.25 | - |
| avg mut/change | 1.000 | 1.593 | 1.485 |
| avg mut/step | 1.000 | 1.731 | 1.000 |
| avg eval count | 436 | 561 | 917 |
| max eval count | 1,556 | 2,359 | 4,134 |
| min eval count | 46 | 48 | 28 |
| fail ratio | 0.000 | 0.000 | 0.000 |

The ranking of the algorithm is the same as for the other inputs with a similar preference of low mutation rates. The RLS has the best performance closely follow by $pmut_{3.25}$ and lastly be the standard (1+1) EA.

| | | | | | | | | |
|--------------------|-----|-----|-----|-----|------|------|-------|-------|
| fails | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS | 997 | 945 | 736 | 51 | 2 | 0 | 0 | 0 |
| RLS_2^S | 972 | 625 | 211 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA (1/ n) | 868 | 436 | 146 | 0 | 0 | 0 | 0 | 0 |
| (1+1) EA (2/ n) | 757 | 277 | 81 | 0 | 0 | 0 | 0 | 0 |
| pmut (3.0) | 781 | 364 | 120 | 0 | 0 | 0 | 0 | 0 |
| pmut (3.25) | 787 | 417 | 128 | 1 | 0 | 0 | 0 | 0 |

No unexpected results for the different input sizes. The RLS variants perform the worst for $n \leq 100$. The input is also rather hard to solve for $n \leq 50$ but not as hard as the overlapped distributed input for example.

| | | | | | | | | |
|--------------------|-------|-------|------|------|------|------|-------|-------|
| avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS | 60 | 169 | 261 | 425 | 390 | 404 | 436 | 570 |
| RLS_2^S | 412 | 1889 | 3552 | 791 | 626 | 591 | 607 | 695 |
| (1+1) EA (1/ n) | 19040 | 15992 | 9758 | 1463 | 883 | 874 | 914 | 1105 |
| (1+1) EA (2/ n) | 20169 | 16391 | 8562 | 1536 | 1088 | 1012 | 1018 | 1076 |
| pmut (3.0) | 26596 | 17183 | 8399 | 1060 | 606 | 578 | 578 | 691 |
| pmut (3.25) | 25871 | 16012 | 7669 | 1229 | 593 | 547 | 553 | 673 |

The input gets easier to solve up until $n = 5000$ and from then on gets harder again with increasing input size. The increase for larger input sizes is much smaller than the decrease for the small values.

| | | | | | | | | |
|--------------------|-------|-------|-------|------|------|------|-------|-------|
| total avg | 20 | 50 | 100 | 500 | 1000 | 5000 | 10000 | 50000 |
| RLS | 99700 | 94509 | 73669 | 5503 | 589 | 404 | 436 | 570 |
| RLS_2^S | 97211 | 63208 | 23902 | 791 | 626 | 591 | 607 | 695 |
| (1+1) EA (1/ n) | 89313 | 52619 | 22933 | 1463 | 883 | 874 | 914 | 1105 |
| (1+1) EA (2/ n) | 80601 | 39550 | 15969 | 1536 | 1088 | 1012 | 1018 | 1076 |
| pmut (3.0) | 83924 | 47328 | 19391 | 1060 | 606 | 578 | 578 | 691 |
| pmut (3.25) | 84210 | 51035 | 19488 | 1328 | 593 | 547 | 553 | 673 |

Mixed and overlapped inputs are best solved by the (1+1) EA with $p_m = 2/n$ for $n \leq 100$ and also relatively good for $n = 500$. After $n \geq 1000$ the standard RLS becomes the best option and it seems like it stays that way for the remaining input sizes.

5.12. Conclusion of empirical results

Bibliography

- [BCP01] Christian Borgs, Jennifer Chayes, and Boris Pittel. Phase transition and finite-size scaling for the integer partitioning problem. *Random Structures & Algorithms*, 19(3-4):247–288, 2001.
- [Dev06] Luc Devroye. Nonuniform random variate generation. *Handbooks in operations research and management science*, 13:83–121, 2006.
- [Die05] Volker Diekert. *STACS 2005: 22nd Annual Symposium on Theoretical Aspects of Computer Science, Stuttgart, Germany, February 24-26, 2004, Proceedings*, volume 3404. Springer Science & Business Media, 2005.
- [DLMN17] Benjamin Doerr, Huu Phuoc Le, Régis Makhmara, and Ta Duy Nguyen. Fast genetic algorithms. In *Proceedings of the genetic and evolutionary computation conference*, pages 777–784, 2017.
- [FGQW18] Tobias Friedrich, Andreas Göbel, Francesco Quinzan, and Markus Wagner. Evolutionary algorithms and submodular functions: Benefits of heavy-tailed mutations. *arXiv preprint arXiv:1805.10902*, 2018.
- [Wit13] Carsten Witt. Tight bounds on the optimization time of a randomized search heuristic on linear functions. *Combinatorics, Probability and Computing*, 22(2):294–318, 2013.
- [Wit14] Carsten Witt. Fitness levels with tail bounds for the analysis of randomized search heuristics. *Information Processing Letters*, 114(1-2):38–41, 2014.

Appendix

A. Appendix Section 1

ein Bild

Figure A.1.: A figure